



Figure 3: RethinkDB interface

- Supports SSL access to have secured access to RethinkDB via public Internet.
 - *More functions*: Supports various mathematical operators like floor, ceil and round.
- Official website:** <https://rethinkdb.com/>
Latest version: 2.3.5

SQLite

SQLite is an open source and embedded relational database, which is designed to provide an easy way for applications to manage data without the overhead. It is highly portable, easy to use, compact, efficient and reliable.

SQLite is ACID-compliant; it implements most SQL standards, and uses dynamically and weakly typed SQL syntax. SQLite engine is not a standalone process like other databases; it can link to static as well as dynamic applications.

Features

- Doesn't require a separate server process or system to operate, and can operate in a serverless environment.
- No requirement for any system administration, and needs a low-configuration machine for build up.
- Self-contained and has no external dependencies.
- Written in ANSI-C, and provides easy and simple API.
- *Cross-platform*: Compatible with UNIX, LINUX, Windows, MAC-OS x, etc.
- Transactions are fully ACID compatible, allowing safe access from multiple processes.
- Supports all SQL queries found in SQL92.
- Fully tested and verified code in SQLite, which is error-free and always up-to-date.

Official website: <https://www.sqlite.org>

Latest version: 3.17.0

Apache Cassandra

Apache Cassandra is regarded as a highly scalable and distributed open source database for managing voluminous amounts of structured data across many commodity servers.

As compared to other open source databases, Cassandra offers various additional high performance capabilities in terms of availability, linear scale performance, simplicity and easy distribution of data across multiple database servers.

Cassandra was developed by Facebook with the prime motive of facilitating Inbox search and was made open source in 2008. It implements the 'Dynamo-style replication model' with no single point of failure, and adds a more powerful 'column family' data model.

Features

- *Massively scalable architecture*: Cassandra has a masterless design, where all nodes are at the same level, which provides operational simplicity and easy scale out.
- *Masterless architecture*: Data can be written and read on any node.
- *Linear scale performance*: As more nodes are added, the performance of Cassandra increases.
- *Fault detection and recovery*: Failed nodes can easily be restored and recovered.
- *Flexible and dynamic data model*: Supports datatypes with fast writes and reads.
- *Data protection*: Data is protected with commit log design and built-in security like backup and restore mechanisms.
- *Tunable data consistency*: Support for strong data consistency across distributed architecture.
- *Multi-data centre replication*: Cassandra provides features to replicate data across multiple data centres.
- *Data compression*: Cassandra can compress up to 80 per cent data without any overhead.
- *Cassandra query language*: Cassandra provides a query language that is similar to SQL language. This makes it very easy for developers moving from a relational database to Cassandra, to use it.

Official website: <http://cassandra.apache.org>

Latest version: 3.10 

References

- [1] <https://www.influxdata.com/>
- [2] <https://crate.io>
- [3] <http://basho.com/products/riak-ts/>
- [4] <https://www.mongodb.com/>
- [5] <https://rethinkdb.com/>
- [6] <https://www.sqlite.org>
- [7] <http://cassandra.apache.org>

By: Prof. Anand Nayyar

The author is an assistant professor in the department of computer applications and IT at KCL Institute of Management and Technology, Jalandhar, Punjab. He loves to work and research on open source technologies, cloud computing, sensor networks, hacking and network security. He can be reached at anand_nayyar@yahoo.co.in. Watch his YouTube videos at [YouTube.com/anandnayyar](https://www.youtube.com/anandnayyar).