

CODE SPORT



Sandya Mannarswamy

In this month's column, we discuss the problem of classifying insincere questions on a community question-answering site.

We have been discussing different natural language processing (NLP) related tasks over the past few months. One of the topics covered was content moderation on social media. Given the explosion of user-generated content on the Web, moderating the content that is shared publicly on online platforms becomes mandatory. Content that is objectionable, hateful and illegal needs to be moderated and filtered. This includes hate speech that incites violence, trolling, spam, and offensive/objectionable content.

Moderating content on social media platforms such as Twitter, Facebook and YouTube has been the focus of much attention over the past few months. Typically, these platforms employ an army of human moderators to scan the content and take down content that is not suitable. However, given the exponential rate of growth of user-generated content, manual content moderation cannot scale. Hence automated approaches that can assist human moderators have become the focus of many research efforts.

There are two major aspects to content moderation. These are: (1) detecting misbehaviour, and (2) detecting misinformation. Misbehaviour includes spam, trolling, scams, fraud, bots, coordinated attacks, cyber bullying, propaganda, extremism and hate speech. Misinformation includes fake reviews, fake news, rumours, hoaxes and fabricated images that spread incorrect information. There are a vast number of channels for misbehaviour and misinformation online. On the corporate intranet, these can include all forms of computer mediated communication including

person-to-person communication such as emails, chats, etc, as well as group communication such as community forums, blog posts, corporate announcements, etc.

On the Internet, content moderation is required on:

- Social media platforms such as Twitter, Facebook and YouTube
- E-commerce platforms such as Groupon, Amazon, etc
- Collaborative and knowledge-based platforms (like wikis) and question-answer platforms (like Quora, StackOverflow, etc)
- Online gaming and transactions platforms

While we have earlier discussed how to identify and filter offensive content and hate speech on social media platforms, in this month's column, we will focus on content moderation for community question-answering platforms.

The prevalence of misbehaviour on focused community question-answering platforms such as StackExchange tends to be less since these forums are typically task oriented and meant for sharing knowledge on specific topics. This enables the easy removal of any questions not relevant to the subject being discussed on the forum. On the other hand, general purpose community question-answering forums such as Quora (www.quora.com) allow questions on any topic. This makes them more susceptible to misbehaviour, since it is not possible to filter irrelevant and insincere questions based on the topic alone.

While most of the questions on such general-purpose CQA (community question and answer)