


# Efficient Log Management with Fluent Bit



According to the official documentation, Fluent Bit is an open source, multi-platform log processor and forwarder that allows a user to collect data and logs from different sources, unify them, and then send them to multiple destinations.

**L**et us explore Fluent Bit, an open source specialised data collector. Apart from covering its use cases, readers will learn how to set up Fluent Bit for a service and write their own plugin.

In recent times, DevOps has become an important part of the software development world. Monitoring and collecting logs is crucial while developing any software. It helps the DevOps teams to keep track of all the events that happened in the past, catch bugs and use the log data to create a better product or service. To keep track of logs in our distributed system, we use log

monitoring. We aggregate the logs and then there are various software that analyse the data.

Logs can be monitored manually, but in a large distributed system there might be millions of log records every day. Manual aggregation, processing and monitoring wastes a lot of time and labour. Hence, we need an automated system for this. This is where Fluent Bit comes in.

If you are looking for a fast log processor that uses less CPU and memory resources, has great throughput and high reliability, Fluent Bit might be your best bet.

Fluent Bit is open source and is written in the C programming language. It is really good for distributed environments where the major considerations are to reduce overheads since the systems have limited capacity. Apart from this, it also provides security for network operations using Transport Layer Security (TLS). It was developed to be cloud-native and has extensive support for Docker and Kubernetes. It is also capable of being plugged and has many input, filter and output plugins. The input plugins are dummy, tcp, syslog, systemd and mqtt, among many others. There are also many

