


Setting Up and Optimising Apache Cassandra

In this article, which is aimed at developers and NoSQL users, we present a step-by-step tutorial on the setting up and optimisation of Apache Cassandra.



Apache Cassandra is an open source, highly scalable NoSQL database, that's well known for handling large volumes of data across multiple data centres and the cloud.

In this article, we will learn to set up and optimise Apache Cassandra.

Setting up Apache Cassandra on Linux

To set up Cassandra on a Linux machine, install Apache Cassandra 3.x on RHEL/CentOS 6.5 or later.

The prerequisites are:

- Yum package management
- Root or non-root user with *sudo* privileges
- Oracle Java Platform, Standard Edition 8 (JDK) or OpenJDK 8. JDK is recommended for Cassandra installation as it has some tools that are not in JRE
- Python 2.7

Installation steps

- Check the Java version, as follows:

```
$ java -version
```

If Oracle Java is used, the results should look like what's shown below:

```
java version "1.8.0_65"
Java(TM) SE Runtime Environment (build 1.8.0_65-b17)
Java HotSpot(TM) 64-Bit Server VM (build 25.65-b01,
mixed mode)
```

- Create a new Yum repo *datastax.repo* under */etc/yum.repos.d*

```
/etc/yum.repos.d/datastax.repo
[datastax-ddc]
name = DataStax Repo for Apache Cassandra
baseurl = http://rpm.datastax.com/datastax-
ddc/3.2
enabled = 1
gpgcheck = 0
```