

“CentOS Linux is built on a lot of past experience”

Q How did you start your journey with the CentOS project?

It was in late 2004. I was not one of the founders of CentOS but showed up on the scene in its early days. At that time, we had a small team and a lot of machines running Red Hat Linux 7.3 and Red Hat Linux 9.

With Red Hat moving down the path towards Red Hat Enterprise Linux, a model that didn't work well for us, I started looking at options. We explored Debian and SUSE initially, but found management and lifecycle on each of them hard to map out workflow into.

It was during this time that I came across the Whitebox Linux effort and then the CentOS Project. Both had the same goal, but the CentOS team was more inclusive and seemed more focused on its goals. So, in late September 2004, I joined the CentOS IRC channel and then, in November, I joined the CentOS mailing list as a contributor. And I am still contributing 13 years down the road!

Q What were the biggest roadblocks that emerged initially while designing CentOS for the community, and how did its core development team overcome them?

A lot of our problems were about not getting off the ground. Initially, there was no clear aim. And then, we faced the challenge that the build systems and code audit tools in 2003/2004 were either primitive, absent entirely or the contributors were unaware of them. A Linux distribution is a large collection

Managing a Linux distribution for a long time requires immense community effort. But what is the key to success in a market that includes hundreds of competitive options?

Also, what are the challenges in building a brand around an open source offering?

Karanbir Singh, project leader, CentOS,

answers these questions and outlines the future of the platform that has been leading Web developments, in an exclusive conversation with **Jagmeet Singh of OSFY.**

of a lot of code, written in many languages — each with its own licence, build process and management.

Three main strategies saw us past that painful process. The first was consistency. Whatever we did, we had to be consistent and uniform across the entire distribution, and make sure all developers had a uniform understanding of the process and flow. The second was a self-use focus. Regardless of what the other people were targeting, all developers were encouraged to focus on their own use cases and their personal goals. The third was the hardest, to try and disconnect commercial interests from developer and contributor work.

Q Why was there a need for CentOS Linux when Fedora and Red Hat Enterprise Linux already existed in the market?

The Fedora project was still getting sorted out around then. Its team had a clear mandate to try and build an upstream-friendly Linux distribution that was going to move fast and help the overall Linux ecosystem mature. Red Hat Enterprise Linux, on the other hand, has been built for the commercial medium to large organisations, looking for value above the code. This left a clear gap in the ecosystem for a community-centric, manageable, predictable enough Linux distribution that the community itself, small vendors, and niche users around the mainstream could consume.

Initially, the work we did was quite focused on the specific use cases that the developers and contributors had. All of us were doing specific things, in specific ways and CentOS Linux fitted in well. But as we started to mature, we saw great success in specific verticals, starting from academia and education institutions to Web hosting, VoIP (Voice over Internet Protocol) and HPC (high performance computing).

Q What were the major outcomes of the Red Hat tie-up?

Red Hat came on as a major sponsor for the CentOS Project in January 2014. From the CentOS Project's perspective, this meant we were then able to start looking beyond just the platform and the Linux distribution. It allowed us to build the infrastructure and resources needed to support