


Sarvam releases 30B and 105B LLMs under Apache 2.0



Indian AI startup Sarvam has released two foundational multilingual large language models (LLMs) as open source software under the Apache 2.0 licence, positioning them as India-built alternatives to global AI systems. The models—featuring 30 billion and 105 billion parameters—are available for public download through AIKosh and Hugging Face, and can also be accessed via Sarvam’s Indus AI chatbot and the company’s API developer dashboard.

The models were first unveiled at the India-AI Impact Summit 2026 in New Delhi, with the official rollout announced on March 6. Sarvam has increasingly emerged as a key contributor to India’s ‘sovereign AI’ initiative, which aims to reduce dependence on foreign AI providers such as OpenAI and Anthropic while developing AI tailored for Indian languages and local applications.

Training was supported by the IndiaAI Mission, with infrastructure from data centre operator Yotta and technical support from NVIDIA, using government-backed GPU compute resources.

Both models use a Mixture-of-Experts (MoE) transformer architecture, activating only a fraction of parameters at a time to improve efficiency and reduce compute costs. The Sarvam 30B model features a 32,000-token context window for conversational applications, while the 105B model offers a 128,000-token context window designed for complex reasoning and agentic workflows.

The models incorporate optimisations such as Grouped Query Attention (GQA) in the 30B model and Multi-head Latent Attention (MLA) in the 105B model to reduce memory requirements during inference. Benchmark tests show the 105B model competing with pt-oss 120B and Qwen3-Next, while outperforming DeepSeek R1, Gemini 2.5 Flash, and o4-mini on the Tau 2 Bench for agentic reasoning and task completion.

“Building these models required developing end-to-end capability across data, training, inference, and product deployment. With that foundation in place, we are ready to scale to significantly larger and more capable models, including models specialised for coding, agentic, and multimodal conversational tasks,” Sarvam said in a blog post.

India should build on existing open source models to power AI: Red Hat

India does not need to build its own foundation or large language models to become an artificial intelligence powerhouse. Instead, leadership will come from how effectively the country adopts, adapts and scales existing systems, according to Navtez Bal, Vice President and General Manager, India and South Asia at Red Hat.

“India will become one of the biggest users of AI in the world. There is no doubt about it,” Bal said. He argued that India’s strength lies in refining available models rather than creating them from scratch. “India will do a much better job in taking the large language models which exist and actually tweaking and distilling them for our own usage, so that we have more fit-for-purpose models.”

Long-term differentiation, he added, will depend on execution. “It’ll be on who uses models more effectively. That’s where the differentiation will lie.”

As sovereign AI gains traction, Bal warned against dependence on proprietary ecosystems. Open source platforms, he said, provide the control governments and enterprises need. “If you are a sovereign AI-focused company, then you need a solution that belongs to you and is not locked in... data remains within India, controls remain within India, the control on source code remains within India.”

Red Hat, built on Linux and OpenShift, positions itself as an alternative to hyperscaler-led stacks and counts enterprises across sectors among its users.

