

learn from extensive general datasets to develop abilities that enable them to identify patterns, comprehend languages and acquire fundamental knowledge about the world, which takes children several years to acquire through reading. It demands substantial resources because it produces basic models that can execute multiple functions.

Users can generate model outputs through different input methods, which include zero-shot, few-shot and role-based inputs, because prompt engineering enables this process without requiring any software modifications. The system needs password access to execute its operations; it manages to complete individual tasks but struggles to process multiple responsibilities that need memory retention of earlier tasks.

Fine-tuning creates a new process which retrains models through targeted data while maintaining permanent weight updates that enable models to achieve deeper specialisation. Through adaptable interface options it delivers stable high-accuracy results, which surpass prompt engineering performance in coding and diagnostic domains.

Fine-tuning can be conducted on regular computer setups and helps SLMs work better than LLMs that use prompt-based systems for specific tasks. The pretraining process establishes the basic system, which temporarily uses prompted guidance to direct its functions until the user establishes permanent system modifications through fine-tuning. The latter renders SLMs highly flexible and efficient.

SLMs in action

Enterprises: Efficiency meets privacy:

Large organisations are rapidly moving mission-critical workloads to fine-tuned SLMs because they can run on-premises or in private clouds. Financial institutions use between 3 and 7 billion models to identify fraud patterns through real-time analysis of millions of transactions that occur at sub-second speed without any data leaving their

security perimeter. Retail and telecom companies use internal customer interaction logs to develop advanced chat and voice agents, which enable them to reduce cloud API expenses by 70 to 90 percent while satisfying their data residency as well as GDPR and CCPA compliance requirements.

Education: Personalised learning

at scale: Educational institutions provide students with interactive learning resources through SLMs that function without the need for internet access or costly subscription services. Students receive personalised reading materials, maths problems and quizzes, which include explanations of content through fine-tuned models that match their specific language needs and learning requirements. TSLMs (time-series language models), which operate on affordable tablets and shared school laptops deliver uninterrupted tutoring services to remote areas of rural India, sub-Saharan Africa and remote Pacific islands as they can function without internet access. Global MOOC platforms use them to auto-generate multilingual subtitles, discussion prompts, and feedback at scale. The result: education becomes more inclusive, equitable, and scalable.

Healthcare: Intelligence at the

point of care: Domain-adapted SLMs operate in hospitals, clinics and remote health posts to process patient notes, lab results and wearable data at local sites. The system generates discharge summaries, which include medication interaction warnings, differential diagnosis recommendations and patient-friendly report translations without transmitting secure PHI data to external servers. Community health workers in under-served areas with slow internet connections use SLMs on tablets or smartphones to assist them with patient triage, maternal care and chronic

disease monitoring. The findings from early 2026 pilots demonstrate that rural areas now experience better diagnostic accuracy and speedier medical treatment because of these advancements.

Research: Democratising

scientific discovery: Researchers from academic institutions and independent organisations can now utilise advanced research tools that do not require them to request GPU access from major technology companies. They can conduct overnight fine-tuning of open models, which include Phi-4 Gemma-3 and Qwen variants, by using domain datasets that contain protein folding sequences, climate model outputs, historical archives and endangered-language corpora. Researchers from Global South universities are now conducting large-scale simulations, literature reviews and data analysis, which computational resources prevented them from doing before. The open source nature of most SLMs fuels global collaboration: a biologist in São Paulo can build on work done by a physicist in Nairobi using the same base model.

All four domains use SLMs to deliver specialised intelligence on local devices, edge servers and private infrastructure. The system provides users with privacy protection, low latency, cost savings, and operating freedom from external service providers.

SLMs are today establishing themselves as a highly effective AI solution. New models like Phi-4, Gemma 3 and Qwen3, along with knowledge distillation methods, enable these models to match large language model performance while maintaining faster, cheaper and more private system advantages. This new era of democratisation enables anyone to access advanced intelligence through basic hardware and original ideas. **END** 🐧

By: Megha Kumari Agarwal

The author works at Tier 1 MNC in Kolkata. She has expertise in IT technologies and database systems, and is also a professional content writer, known for simplifying complex tech topics for a wider audience.