Devin, OpenAI's coding tools, and Anthropic's agents to further streamline development.

Q. What are the biggest scaling challenges in today's AI, and how does your tech fix them?

A. The core challenge with genAI today is scalability—it simply isn't there yet. Frameworks like vLLM are difficult to scale, and only a few experts know how to manage tens of thousands of concurrent requests with proper autoscaling and parallelism across nodes. Optimising tensor and replica parallelism based on specific hardware is complex and still being figured out. Companies like OpenAI, Anthropic, and Google can scale because they focus on a small set of known models and build custom optimisations. But creating a general-purpose inference runtime that works across any model, hardware, and cloud setup is far more difficult.

Q. Is your software runtime open source and built from scratch?

A. Yes, a part of our software is open source. Our base system was built from scratch, but it does include some open source parts like many other projects. We've open sourced part of the project—around 20–30%—mainly some tools and utilities we used to help build our system. It's not fully built on top of another open source project. We're also starting to grow a community around it and planning to do more outreach soon.

Q. Are the developers from the open source community or part of your team?

A. Right now, all the development is done by our team and a few partner groups. But in the next few months, we want to grow the community and get more outside developers involved. The code is already available on GitHub.

Jithin VG receiving the Breakthrough Innovation ISV award from Intel Corporation

Q. What are the advantages of open sourcing a genAI project?

A. We believe this technology can make generative AI far more affordable and accessible. Today, running large models like LLaMA 8B requires expensive GPUs—often costing tens of thousands of dollars a month—which most small and mid-sized businesses simply can't afford. While APIs from providers like OpenAI or Anthropic seem cheaper initially, costs quickly rise with scale, and users remain dependent on those platforms without owning any part of the technology.

That's why open source genAI is crucial. It empowers developers and companies of all sizes to build, run, and improve AI systems on their own terms. An open ecosystem drives innovation, enables problem-solving across diverse use cases, and makes the technology more transparent and trustworthy. With more people contributing, we can push genAI to be faster, more efficient, and widely useful—beyond just those with deep pockets.

Q. What challenges do you face when open sourcing genAI systems?

A. A key challenge with open sourcing is building an active community early on. You need enough contributors to help the project grow and scale. Another issue, especially in India, is the shortage of developers skilled in low-level systems work—like firmware, kernel writing, or programming in C and Assembly. These experts are hard to find, which makes early development difficult.

India's developer community is great at using open source tools, but not as strong in contributing back. We consume a lot from GitHub, possibly more than any other country, but we produce fewer major open source projects. Growing a culture of contribution is essential, and while it's improving, it's still a major challenge—especially for projects being built from here.

Q. Are there any parts of your tech stack that you haven't open sourced yet, and why?

A. There's a small part of our stack—a new runtime for non-LLM use cases—that isn't open sourced yet. We're still deciding whether to release it and, if so, which parts. Our team is in the middle of that decision-making process.

While we're a research-focused company, we also need to generate revenue to stay afloat. So, for now, we're keeping some IP in-house to work directly with large enterprises instead of giving everything away on GitHub. Our tech team strongly supports open source and would prefer to release everything from day one, but we're taking a case-by-case