approach. Most of our work will likely be open sourced over time, though we don't yet have a fixed process for when or how that will happen.

Q. What open source licence do you use for genAI runtime? How does it affect business use?

A. We use three licences: MIT, Apache 2.0, and our own Bud Community Licence. The Bud licence protects us by stopping companies from just taking the code to offer it as a paid service, since we plan to launch our own SaaS soon. It also limits use by companies valued over $100 million, encouraging them to get an enterprise licence.

Licensing is a balance. Too strict, and fewer people use it. Too open, like MIT or Apache, and it can hurt future business opportunities. For example, LLaMA's community licence is strict but still widely used. In generative AI, strict licences aren't a big issue. The key is to pick a licence that fits your goals.

Q. How do you add new features without breaking old ones in open source AI?

A. We've prioritised rapid feature development while ensuring backward compatibility. From the start, the codebase was carefully designed with long-term stability in mind—not a quick hack, but a well-planned architecture inspired by mature industry projects like OpenVINO, OpenCL, and ONNX. By studying their architectures and communities, we identified common challenges—especially those causing repeated structural changes and backward incompatibility—and worked to avoid them. Although future requirements may still force major changes, our planning so far has helped maintain stability across multiple iterations.

That said, we're not focused on releasing features rapidly for the sake of speed. As an optimisation system, our priority is building reliable, high-performance components. We take our time to ensure every line of code is efficient and stable, emphasising quality over quantity. Our goal is to deliver lasting value and performance rather than frequent feature updates.

Q. How do you handle security risks in open source AI runtime code?

A. From the start, we had a clear vision — to build an open source project focused on enterprise needs, especially secure on-premises deployments. Security, along with performance and scalability, has been our top priority, shaping the entire system architecture.

We follow key security standards like ISO, MITRE, CWE, SOC, HIPAA, and AI-specific guidelines from the White House and others. Built for compliance from day one, we also use AI tools to monitor vulnerabilities. Our proprietary BUD SENTRY or Secure Evaluation and Runtime Trust for Your models, is a zero-trust model ingestion framework that ensures that no untrusted models enter your production pipeline without rigorous checks.

Q. How do you see India and Indian developers being part of this?

A. India has one of the largest developer communities in the world, so it will play a key role in building and supporting our project. The country also needs a solution like this right now. We're working to catch up with nations like the US, China, and those in the EU in terms of AI and data centre capabilities, but India still lags behind—while the US has around 21 gigawatts of data centre capacity, India has only about 1 gigawatt. One way to close this gap is by using existing infrastructure more efficiently—like CPUs or commodity hardware —instead of relying solely on costly GPUs. Our software enables this, helping developing countries make faster progress in AI with the hardware they already have. That's why we believe this will be especially valuable for India and why we expect strong community support from Indian developers.

Q. How will genAI hardware and software evolve in five years?

A. We ask this daily because it drives us. Since genAI aims to automate everything, we asked: shouldn't it automate itself? That led to Bud Agent—an AI that builds and manages genAI infrastructure, rewrites its code, and acquires hardware as needed. It's not open source yet, but videos are available on LinkedIn. We believe genAI will soon run itself, optimising and scaling without humans.

GenAI will evolve towards general and maybe superintelligence. These terms are unclear since we don't fully grasp intelligence yet. But I hope within a few years we'll create AI as capable as humans and start solving these big questions.

Q. What are your plans and what are you working on now?

A. India lacks fundamental AI research and mostly consumes technology rather than creates it. Most genAI startups rely on models and architectures developed outside the country—mainly in the US, China, and Europe. Less than 0.01% of practical genAI research papers come from India. While regional personalisation is important and many Indian companies focus on it, this doesn't address the deeper gap in core AI research and innovation.

Our goal is to change this by focusing on fundamental AI research, creating new, efficient model architectures tailored for India and similar developing countries. These models would be optimised for low energy, low compute, and edge device deployment. Rather than just building huge data centres like the US, India should leverage its billions of existing client devices to maximise compute capacity and compete globally. Innovation and deep research are our priorities because we believe growth and success naturally follow when you create real value. END