

Possible enhancements:

- Add patient history and lab report embeddings to refine suggestions.
- Integrate with Indian EHR systems (eSanjeevani, Ayushman Bharat Digital Mission).
- Run entirely on low-cost local servers or edge devices (offline).

pgvector vs other vector databases

As vector search becomes mainstream in AI and data applications, a growing number of vector databases (e.g., Pinecone, Weaviate, Qdrant, Milvus) are entering the market. However, the cost structure and integration complexity of these solutions vary greatly, especially when deployed in real-world enterprise scenarios (Table 3).

Table 3: A comparison of pgvector with other vector databases

Feature/ Cost	pgvector (Postgres)	Pinecone/ Weaviate (SaaS)	Faiss (self-hosted)
Setup cost	Minimal (Postgres extension)	High (requires account and integration)	Medium (needs hosting + config)
Hosting	On-prem/ Cloud/ Docker	Cloud only	On-prem or cloud
Vector + SQL integration	Seamless	Limited or non-existent	Requires custom dev work
Security	Full data control	Vendor-managed	Fully controlled
Multi-modal search	Hybrid (SQL + vectors)	Mostly unstructured	Complex to integrate
Annual cost (mid-scale)	Rs 50,000–150,000 (\$700–1800)	Rs 600,000–1.2 million (\$7000–15,000)	Rs 200,000–350,000 (\$2500–4500)

Here's a key insight. If you're already using PostgreSQL, adopting pgvector is almost cost-free in terms of setup and maintenance. You gain vector capabilities without sacrificing relational power or requiring new infrastructure.

The future of vector search in hybrid SQL databases

As AI continues to evolve, the convergence of semantic intelligence and structured querying is redefining how data systems are designed. This has led to the rise of hybrid SQL

+ vector databases, a category in which PostgreSQL with pgvector is emerging as a front-runner.

Most real-world data systems involve both structured filters (e.g., WHERE region = 'South India' AND year > 2022) and unstructured, semantic matching (e.g., "similar customer complaints", "related incident reports"). Traditional vector databases lack SQL capabilities, making it hard to express complex business logic. On the other hand, traditional SQL databases couldn't handle embeddings, until now (Table 4).

Table 4: Evolution roadmap

Phase	Capability	Description
1	Standalone vector DBs	Fast, purpose-built, but limited to unstructured data
2	SQL + Vector search (pgvector)	Combines structured filters + semantic similarity
3	Native ML/AI extensions in SQL engines	Predictive modelling and RAG flows directly inside SQL
4	LLM-native databases	Query understanding + natural language interfaces
5	Autonomous AI-driven query optimizers	Real-time tuning, vector-aware planners, intent prediction

Hybrid vector databases represent a crucial step forward—bridging the gap between relational logic and semantic understanding. By combining the maturity and reliability of PostgreSQL with the flexibility of vector search, developers can now build intelligent, real-time applications that reason, recommend, and respond like never before. Whether it's powering a search engine that understands intent, detecting engine anomalies in aviation, or assisting doctors in remote healthcare centres, hybrid databases enable AI-native workflows within trusted ecosystems.

More importantly, pgvector allows developers to embrace these capabilities without abandoning the tools, skills, or infrastructure they already use. If your next application needs to search smarter, recommend faster, or reason deeper, PostgreSQL with pgvector might be the database you need.

The future is not about replacing SQL with AI — it's about embedding AI into SQL, responsibly, scalably, and intelligently. **END** 🐧

By: Praveen Kumar Magesh

The author is a cloud database engineering professional with over 18 years of experience in architecting and optimising large-scale database systems across global enterprises. His focus areas include PostgreSQL, Oracle, cloud-native database platforms and automation.