

IBM launches Open Source Cloud Guide for developers

IBM has released the Open Source Cloud Guide, which highlights various use cases that are important in hybrid cloud environments. It features the important open source projects in these areas, and discusses how various clouds are using open source in their offerings.



“As both the cloud and open source landscapes evolve, we saw a need for a guide to highlight important aspects of hybrid cloud and multi-cloud development—and their corresponding open source services,” IBM said.

According to an O’Reilly survey that IBM commissioned in 2020, developers who want to build cloud applications should work on honing their open source skills instead of only focusing on developing skills for a proprietary cloud. Because every major cloud platform uses open source software in its infrastructure, developing skills related to open technology makes developers more desirable to potential employers and helps developers compete in hybrid environments.

Todd Moore, vice president, open tech, and Chris Ferris, CTO, open tech from IBM noted that the survey highlighted that the most desired of the open source skills are around Linux (containers), artificial intelligence and machine learning, and data storage.

“This got us thinking: How do those skills translate exactly to developing for hybrid cloud environments, inclusive of the major cloud providers? What open source technologies are most used? The purpose of the Open Source Cloud Guide is to answer those questions,” reads the IBM developer blog.

Developers can discover more and contribute to the guide on GitHub.

The funds are grouped in three: foundation memberships (52 per cent), long-tail projects through GitHub Sponsors and Open Collective (36 per cent), and internships for new contributors through Outreachy (13 per cent).

“All tech companies stand on the shoulders of community-supported open source giants, and Sentry is no exception. With this fund we prioritised support for our dependencies in order to strengthen our supply chain. But, more than that—Sentry itself was a volunteer-run project for many years. Yes, we took a commercial route, but we respect the many projects that have chosen a different path. Maintainers should be able to determine their own future, and financially supporting our community-managed dependencies makes that a bit more feasible for them,” wrote Chad Whitacre, senior software engineer at Sentry in a blog post.

The company has audited its product architecture and generated a list of seven major community-led open source projects including Python, Django, Rust, JavaScript, PostgreSQL, Apache, and Linux. “These projects are all backed by formal non-profit foundations; we added an eighth foundation, the Open Source Initiative, to represent the open source community as a whole. We decided to allocate half of our budget (52 per cent) to these eight foundations,” said Whitacre.


CloudQuery raises US\$ 3.5M in funding

CloudQuery, a startup giving developers visibility into their cloud infrastructure assets and configuration, has announced that it has closed a US\$ 3.5 million seed funding round led by Boldstart ventures, with participation from Work-Bench, Mango Capital and Haystack.

As cloud infrastructure providers and service catalogues have grown, the burden on developer workflow has been ever more custom scripts, hacking and hassle for developers to get the insights they need over their sprawling cloud infrastructure. Each cloud provider has its own proprietary tooling for infrastructure visibility — AWS Identity and Access Management (IAM), Azure Cloud Discovery, Google

Cloud Asset — each with its own learning curve, and each siloed within that single cloud provider environment.

“With cloud infrastructure, developers too often find themselves blind about what their assets are and what they are managing,” said Yevgeny Pats, co-founder and CEO at CloudQuery. “They only get spotlights on small parts of their infrastructure, and spend a lot of cycles writing manual scripts for specific APIs,



doing transformations, and managing code. The vision for CloudQuery is ‘dev-first’ visibility into infrastructure. That means continuous extract, transform, load (ETL) of your cloud infrastructure assets into a relational database, a simplified SQL query model that lets you ask all the questions you need to know for your use cases, and a vibrant community of contributors that lets you tap into pre-supported cloud providers and queries, so you don’t have to build everything from scratch yourself.”

CloudQuery integrates with the cloud service providers, including AWS, Azure and Google Cloud. It fetches read-only data and normalises that data (via ETL) into a relational database. By exposing cloud infrastructure data as SQL-queryable, CloudQuery does not require mastering new DSLs. And the project boasts more than 100 pre-written queries for all the cloud providers that it supports, so developers enjoy a running start for all the common queries they have of their cloud assets – but CloudQuery makes it easy to build custom queries too.

For more news, visit www.opensourceforu.com

