

Deploying Generative AI Models Efficiently

Enterprise deployment of generative AI hinges on the optimisation of hardware and software. It's important to choose the right-sized language model and fine-tune it as per your need.

OpenAI's launch of ChatGPT powered by GPT-2 in mid-2020 showcased a model with 175 billion parameters—a monumental breakthrough at that time. By the time GPT-4 arrived, parameter counts had surged into trillions, enabling sophisticated chat assistants, code generation, and creative applications, yet imposing unprecedented strain on compute infrastructure.

Organisations are leveraging open source genAI models such as Llama to streamline operations, improve customer

interactions, and empower developers. Choosing a large language model (LLM) optimised for efficiency enables significant savings in inference hardware costs. Let's see how this can be done.

Choosing the right-sized LLM

Since the public launch of ChatGPT, generative AI adoption has skyrocketed, capturing the imagination of consumers and enterprises alike. Its unprecedented accessibility has empowered not just developers but also non-technical users to embed AI into

their everyday workflows.

Central to this evolution is a fundamental measure of progress: LLM parameters — the trainable weights fine-tuned during learning to determine model capability. In 2017, early generative AI models based on the Transformer architecture featured roughly around 65 million such parameters.

This explosive growth has reinforced the belief that 'bigger is better', positioning trillion-parameter models as the benchmark for AI success. However, these massive models



The number of people using genAI