

Apple releases Ferret, an open source multimodal LLM



In an unexpected move, Apple has discreetly unleashed Ferret, a groundbreaking open source multimodal large language model (LLM) created in collaboration with Columbia University. Revealed on GitHub, Ferret challenges the norm with its unique prowess in referring to image regions in any free-form shape and automatically grounding relevant text.

Ferret, curated on the GRIT data set featuring 1.1 million samples, demonstrates superiority in classical referring and grounding tasks. It outshines existing multimodal language models (MLLMs) in region-based and localisation-demanded multimodal chatting, exhibiting improved image detail descriptions and a notable reduction in object hallucination.

Acknowledging the occasional production of counterfactual responses, the researchers confirm plans to elevate Ferret's capabilities by incorporating segmentation masks and bounding boxes.

Apple's decision to release Ferret without ceremony underscores its commitment to leading in multimodal AI. The uncharacteristic open source approach signals a departure from Apple's traditional closed-door strategy, hinting at potential collaborations and community-driven advancements.

Ferret's strength lies in its simplicity, offering more than basic image identification. It seamlessly connects elements within images, promising applications in image search, accessibility, and scenarios requiring nuanced comprehension.

Developers diving into the GitHub repository might unearth transformative possibilities for Ferret, marking a significant leap towards accessible and innovative AI solutions.

Mistral AI to release open source GPT-4 in 2024

Mistral AI, the Paris-based tech trailblazer, announced on French national radio that it will open source its highly anticipated GPT-4 level model slated to be released in 2024. Arthur Mensch, Mistral AI's CEO, conveyed the ambitious move, stating, "We believe in the power of collaboration and the democratisation of AI. By open sourcing GPT-4, we aim to empower developers and researchers worldwide to contribute and benefit from cutting-edge language models."

Mistral AI's recent triumphs include securing a whopping \$415 million in funding, elevating its valuation to an impressive \$2 billion. Leading the charge in the funding round was Andreessen Horowitz (a16z), with Lightspeed Venture Partners reaffirming their commitment.

To address sustainability hurdles faced by open source language model (LLM) firms, Mistral AI unveiled 'La Plateforme', providing API endpoints for its models as a revenue-generating solution. It has categorised its models into Mistral Tiny, Mistral Small, and the upcoming Mistral Medium.



The Mistral Medium model, currently in development, is supposedly proficient in multiple languages and code. It targets top-tier performance, boasting an MT-Bench score of 8.6, outshining the renowned GPT 3.5.