

Responsible AI is the Need of the Hour

With hallucinations, bias, opaque decisions, and CO₂ costs adding up, AI needs discipline and responsibility built in from the start.



Cases of AI causing real-world harm increased noticeably in 2024-25. Responsible AI has now become a mandatory compliance layer for enterprise stability, moving from 'good practice' to a fundamental design condition.

Global incident trackers, including the AI Incident Database and the OICT Monitor, indicate a marked escalation in AI-related harm. Within a matter of months, incident counts increased from approximately 140,000 to more than 150,000, thereby raising the risk share from around 50 per cent to over 60 per cent. The expanding set of risks includes multimillion-dollar deepfake CEO impersonation scams, ransomware attacks enabled by AI tooling, discriminatory outcomes, unsafe

predictions, and privacy breaches. Misuse is now easy, scalable, and cheap.

Even generative video systems show clear, high-visibility bias. In a widely noted example, a model produced only male images when asked about CEOs and only female images when asked about flight attendants. These outcomes are not fringe glitches; they expose structural weaknesses across modern AI systems and mirror what we see in everyday use cases, from loan-approval bots rejecting applicants unfairly to e-commerce customers telling us their AI-first experiences are too costly to scale.

For me, this surge changes the framing entirely. Deepfakes, cyber fraud, biased predictions, privacy breaches, and misuse are operational

risks that are affecting consumers, enterprises, and public institutions.

AI models continue to struggle with hallucinations and inconsistent predictions. Healthcare provides some of the clearest evidence. A model incorrectly predicting diabetes for an individual led to the wrong medication being prescribed, with serious consequences. Another model offered flawed medical suggestions that understandably alarmed users. When systems behave this way, trust erodes faster than technical teams can explain what went wrong.

The environmental impact adds another dimension. A single ChatGPT-style prompt produces about 4.32 grams of CO₂. That may seem minor, but on the enterprise scale it adds up fast. When millions of queries flow through models running in non-optimised data centres, the emissions rise sharply.

Operating costs also directly correlate with model maintenance and can increase with scalability challenges. The need to maintain constant accuracy in the face of drift and other risks further complicates responsible deployment.

Principles of responsible AI

Principles of responsible AI ensure that AI continues to work equitably and does not dilute results based on the different characteristics of users. Here is a quick look at the principles that I recommend.

Ethical usability: AI applications should be created only for ethical usage.

Fairness: This is where it all starts. AI must treat everyone fairly and not