

Violations of the banned-system provisions may result in penalties of up to 6 per cent of global revenue from AI or €35 million. This is one of the strongest reasons I describe responsible AI as a mandatory compliance layer rather than a voluntary commitment. Regulatory expectations are now explicit: responsibility is compulsory.

Industry frameworks and toolkits provide practical guidance for enterprises to help build auditable and fully compliant AI systems.

Gartner's PRISM model is a checklist to verify that each stage of AI development, from data collection through final deployment, adheres to responsible standards. The IEEE 7000 series goes a step further by embedding ethics directly into system design. This can help developers and AI architects follow responsible AI standards from the ground up. Similarly, ISO 42001 is a set of globally recognised guidelines for organisations to adopt consistent responsible AI practices irrespective of their industry or geographical location.

On the practical side, responsible AI toolkits provide hands-on support. Microsoft's toolkit provides open source APIs that enable the integration of fairness, security, and ethical compliance checks directly into AI systems. IBM's Fairness 360 toolkit equips teams with tools to detect bias, mitigate unfair outputs, and ensure transparency in model predictions. Infosys' Responsible AI toolkit provides APIs that validate whether AI applications adhere to ethical and governance principles, maintain fairness, and meet regulatory guardrails.

Tools that operationalise responsibility

Putting responsible AI into practice means being careful and consistent at every stage of an AI system's life. To reduce the chances of unfair or discriminatory outcomes, I rely on a combination of practical methods to




Figure 1: Responsible AI

keep the system trustworthy:

- Diverse and representative datasets
- Fairness-aware algorithms and testing
- Demographic audits
- Bias detection and mitigation tools (IBM Fairness 360, Microsoft Fairlearn, Google What-If)
- Human-in-the-loop reviews
- Regular, structured audits
- Diverse teams that expand perspective and reduce blind spots
- Explainability tools like SHAP and LIME
- Adopt a governance framework tailored to the specific context in which the AI application will be deployed.

Sustainability measures are just as critical in the scheme of AI adoption responsibility. For enterprise-level organisations, small steps lead to significant results. You can optimise models for efficiency, use green data centres, and deploy energy-efficient hardware to reduce overall CO₂ emissions.

Enterprise boards have now recognised the need for responsible AI. AI risks now hold the same level of importance as cybersecurity or financial exposure. Companies now have governance committees to enforce responsible AI standards.

All development and expert opinions converge on a single clear resolution. Responsible AI standards must be incorporated from the outset and should be present at all stages of design, development and deployment. This is the only way to ensure the system remains stable and secure.

Trust grows when users know that AI will behave as expected. A strong commitment to ethical standards also ultimately builds the AI brand's reputation. **END** 🐧

This article is based on the session titled 'Responsible AI in GenAI and LLM Systems' by Selvi Nandakumar, director of product management, Infosys Equinox, at AI DevCon in Bengaluru. It has been transcribed and curated by Apurba Sen, senior journalist at the EFY Group.