

## IndiaAI prioritises strict bias checks on open source LLMs



Entities developing large language models (LLMs) under the IndiaAI Mission have been instructed to make bias mitigation a top priority before deployment. Given India's diverse social landscape, officials from the Ministry of Electronics and Information Technology (MeitY) emphasised that government-supported AI models must avoid producing insensitive or discriminatory outcomes related to caste, gender, food practices, regional, linguistic, ethnic, or religious differences.

A MeitY official said, "Sensitive connotations linked with caste, gender, food practices, regional and linguistic stereotypes, as well as ethnic and religious differences have to be handled with utmost care. We want Indian models to be inclusive, and not discriminatory or based on historical biases. As a result, all the under-construction LLMs have been told to integrate stringent stress testing into their framework."

Another official added, "The sovereign LLMs are a crucial milestone of our AI journey, which should bring the country together. The focus should not shift from this, towards bad actors or miscreants trying to orchestrate a frenzy by feeding compromising prompts to AI. We need to be careful since machine learning tools process data on a massive scale; even small biases in the original training data can lead to widespread discriminatory outcomes."

Startups and consortia developing indigenous AI include Param-1, BharatGen, Tech Mahindra's Indus LLM, Sarvam AI's open source Sarvam M, and Soket AI's DHRITH.

### Kindred Tech develops Katalyst for child safety in India

India is witnessing a sharp surge in online child exploitation, with 2.2 million cybertips received by Indian agencies in 2024, according to the National Centre for Missing and Exploited Children (NCMEC). Addressing this challenge, Kindred Tech, a non-profit organisation, has developed Katalyst, an AI-driven digital investigation platform designed to combat online child exploitation, human trafficking, and other forms of technology-enabled abuse.

### Sonatype Guide brings open source intelligence to AI models in real time

Sonatype Inc. has introduced Sonatype Guide, a security-first system designed to bring real-time open source intelligence into AI-assisted software development. Most AI models are trained on public data that may be months or years old, causing coding assistants to recommend vulnerable, low-quality or even hallucinated packages. A Sonatype study shows that leading generative AI large language models hallucinate packages up to 27% of the time, creating unnecessary rework, slowing delivery, wasting LLM tokens and exposing teams to significant security risks, including malicious open source components.

Enterprises testing Sonatype Guide reported more than 300% improvement in security outcomes, along with major reductions in remediation. Dependency-upgrade costs improved by over five times, measured in both developer hours and direct spend.

Bhagwat Swaroop, chief executive, Sonatype, said: "Guide brings discipline and intelligence to AI-assisted development. It empowers teams to move faster and more safely by steering AI towards secure, reliable components and automating tedious dependency work that slows them down."

Sonatype Guide integrates with tools such as GitHub Copilot, Google Antigravity, Claude Code, Windsurf, IntelliJ with Junie, Kiro from AWS and Cursor. Core features include a model context protocol server for real-time package interception, enhanced open source software search and an enterprise-grade API.

