

openSUSE empowers Linux developers with Intel NPU access

The openSUSE project has begun distributing packaging for the Intel Neural Processing Unit (NPU) driver, enabling small-scale AI development on Linux. Support is available through openSUSE's Tumbleweed rolling-release distribution, giving developers direct access to Intel's built-in CPU accelerator for AI workloads on laptops and PCs.

NPUs offer better performance and lower power consumption than traditional CPUs and GPUs on smaller hardware, making AI more practical for everyday devices. To utilise the NPU fully, developers require OpenVINO user-space software, developed by Intel and optimised for low power, low latency, and edge deployment. This allows AI applications to run on existing hardware without costly upgrades.

The open source kit has already supported the creation of Intel's modular, edge-native software platform, released last year. Meanwhile, the OpenVINO Model Server is designed for cloud and server deployments and remains compatible with hyperscaler cloud storage.

"With the arrival of NPU (Neural Processing Unit) devices, it has become possible to run LLMs and generative AI applications without relying on a dedicated GPU ... As members of an open source community, we have a responsibility to democratize emerging technologies and help reduce the digital divide, and this delivery is another important step in that direction," said Alessandro de Oliveira Faria, innovator, Intel.



Built on open, AI-ready technology principles, Katalyst automates the ingestion, prioritisation, and analysis of cybertips, reducing reliance on slow, manual review processes that can average 35 pages per report.

Katalyst combines advanced case management, a secure media library, AI-driven categorisation of sensitive media, and automated computer forensics support to streamline investigations and enable coordinated, cross-jurisdictional responses. Its architecture is built on MongoDB Atlas, allowing processing of structured and unstructured data at scale, storing AI inferences alongside raw case data, and supporting vector-based similarity searches for victim identification. Offline-ready capabilities ensure field usability in low-connectivity environments.

The platform's impact is already visible. In 2024, Kerala Police's Counter Child Sexual Exploitation Centre (CCSE) completed an 18-month pilot using Katalyst, resulting in 96 arrests, safeguarding 20 children, and issuing 18 international referrals, earning the Technology Sabha Excellence Award 2023.

"Our mission is to empower investigators and law enforcement agencies with tools that surface insights faster and ensure every child's case receives the urgent attention it requires," said Bree Atkinson, CEO, Kindred Tech.

Kindred Tech plans to expand Katalyst across additional Indian states, integrating advanced vector search and predictive analytics to support real-time, cross-border investigations.

The new Red Hat AI 3 integrates natively into Kubernetes

Red Hat has launched Red Hat AI 3, an open source distributed inference engine integrated natively into Kubernetes, providing a unified collaboration layer for developers, operators, and data teams. Key tools such as AI Hub, Gen AI Studio, and Model as a Service enable shared workflows across teams.



Paul Nashawaty, practice lead and principal analyst at theCUBE Research, said, "These combined features transform Red Hat AI 3 into an enterprise collaboration fabric for AI, where developers, operators and data teams share

one control plane, one governance model and one source of truth."

Robert Shaw, director of engineering at Red Hat, noted, "Reasoning capabilities that have been added to some of the proprietary models are starting to show up in the open source models. There is continued improvement in the overall quality of those open source models, which is what actually unlocks the business use cases and consumer use cases that are powering all this frenzy."

Red Hat also focuses on security and zero-trust AI, embedding ongoing verification, access controls, and compartmentalisation to protect AI workloads. Rob Strechay, analyst at theCUBE Research, added, "Red Hat looks to unveil a strategy to secure AI models throughout the entire lifecycle."

AWS taps SUSE to supercharge Amazon Linux

SUSE has announced a major collaboration with Amazon Web Services (AWS) to strengthen the cloud native and open source capabilities of Amazon Linux. The agreement introduces the new Supplementary Packages for Amazon Linux (SPAL) service, delivering thousands of additional enterprise-grade open source packages to customers relying on Amazon Linux for mission-critical operations.

Rick Spencer, GM of business critical Linux at SUSE, said: "We're proud that the Amazon Linux team has chosen to leverage SUSE's Linux expertise and secure