

Mitsubishi Electric is now a Gold member of the Linux Foundation

Mitsubishi Electric Corporation has upgraded its membership in the Linux Foundation from Silver to Gold, marking its evolution from a user to a significant open source contributor.

Jim Zemlin, executive director, Linux Foundation, said: “We are thrilled to have Mitsubishi Electric upgrade to a gold membership with the Linux Foundation. Their established history in nearly every industry, from the edge to automotive solutions, security to manufacturing, underscores the importance of collaboration in software development. We are eager to continue working alongside Mitsubishi Electric on open source, both in the Japanese market and beyond.”

Mitsubishi Electric released the Serendie Design System, a UI/UX framework, in November 2024 to enable co-creation internally and externally. In early 2025, the company established the Open Source & InnerSource Program Office (OSPO/ISPO) to promote open source best practices and bridge internal silos.

Shingo Oidate, general manager, OSPO/ISPO, Mitsubishi Electric Corporation, stated: “Upgrading to a Gold member reflects our firm belief in the importance of open source, and demonstrates our commitment not only as a user, but as an active contributor, expanding our ability to collaborate globally and deliver transformative value to society.”



language- and environment-neutral, supporting Python, JavaScript, Rust, or any file-based workflow, across cloud or on-prem setups.

Tobi Lutke, CEO of Shopify, said: “Tangle is a major piece of our Shopify data and ML system. It makes complex things easy and automatically avoids doing things more than once, saving an insane amount of waste.”

TrueReach AI releases Entropy, its autonomous SDLC platform

TrueReach AI has announced Entropy, an AI-native, autonomous software development life cycle (SDLC) platform that directly exposes the limitations open source AI and DevOps ecosystems continue to face -- fragmentation, orphan code, and the absence of end-to-end orchestration.

Positioned beyond AI-assisted development, Entropy represents a shift towards AI-owned SDLC execution, where humans validate, govern, and ensure quality. The platform replaces fragmented tools with a single autonomous engine spanning architecture design, sprint planning, coding, testing, and validation, delivering production-grade enterprise software rather than isolated code artefacts.

The platform has already been validated in production, autonomously building a full-scale enterprise system deployed across 65 cities, digitising sales and customer operations with zero human-written code. What traditionally required 120 man-weeks was delivered in 24 man-weeks, marking a 75% reduction in development time.

AMD goes open source with ROCm to challenge Nvidia Cuda

AMD is doubling down on open source to make its ROCm platform a go-to ecosystem for AI development, aiming to break Nvidia Cuda’s dominance. The US chipmaker is positioning ROCm as a collaborative software stack, integrating



with projects like virtual large language model (vLLM) to accelerate innovation.

“We could try to build something that’s closed source, but we won’t get the velocity of an open ecosystem. Instead, we want to leverage everyone’s capabilities to move the industry

forward. It’s like the Linux kernel, where everyone collaborates and moves really fast,” said Anush Elangovan, AMD vice-president of AI software.

Targeting the Asia-Pacific region, AMD hopes ROCm will provide a ‘common baseline’ for companies to scale AI capabilities without relying on proprietary tools. Elangovan noted, “I’ve been involved with quite a few companies in the region that are building large data centres and deploying AMD chips at scale. ROCm will also allow them to compete in both model development and infrastructure.”

The ‘ROCm everywhere’ initiative ensures a unified experience across devices, from laptops to supercomputers, supporting students, startups, and enterprise deployments.

ROCm 7, released in September 2025, introduces native support for the AMD Instinct MI350 and MI325X GPUs, lower-precision data formats (FP4, FP8) for faster AI inference, broader Windows and consumer GPU support, and day-zero integration with tools like PyTorch and vLLM.

“You can go for a little less density, so you can do air-cooled infrastructure versus liquid-cooled infrastructure, and then still get the capabilities that are top of the line,” added Elangovan, highlighting scalability advantages.