

Arcee AI launches two Apache-licensed Trinity models

Arcee AI has launched Trinity Mini and Trinity Nano Preview, the first models in its Trinity family, marking a rare US-trained, fully open-weight MoE release under the Apache 2.0 licence. The suite is positioned as a direct response to China's lead in open-weight innovation, offering enterprises and developers unrestricted commercial use, full modification rights, and on-prem deployment without foreign dependencies.



Trinity Mini, a 26B-parameter model with strong results in MMLU (84.95), Math-500 (92.10), GPQA-Diamond (58.55) and BFCL V3 (59.67), is engineered for long-context reasoning and high-throughput tool use. Trinity Nano Preview, at 6B parameters, is a lighter, chat-focused variant. Both employ Arcee's AFMoE architecture with sigmoid routing, gated attention, and depth-scaled normalisation for training stability.

A 420B-parameter model, Trinity Large, is now in training using 20T curated and synthetic tokens, set for release in January 2026. If successful, it will be one of the only US-trained, fully open-weight frontier models.

The models are available on Hugging Face, OpenRouter and chat.arcee.ai, with early integrations across Benchable.ai, Open WebUI and SillyTavern.

Ukraine builds sovereign AI using Google's Gemma

Ukraine has launched a landmark national AI initiative, building a fully sovereign large language model (LLM) using Google's open source Gemma framework. The project aims to give Ukraine complete control over its AI systems, ensuring they serve both military and civilian purposes while avoiding dependency on foreign, closed AI platforms like OpenAI's ChatGPT.



The AI will be trained to handle Ukrainian context, dialects, and minority languages, including Crimean Tatar and regional blends of Ukrainian, Russian, and Bulgarian, addressing the shortcomings of existing models. Officials

emphasise that this localisation will improve accuracy in legal, cultural, and military applications.

Oleksandr Bornyakov, deputy minister of digital transformation, stated, "avoiding reliance on systems like OpenAI's ChatGPT was intentional, as Ukraine's armed forces plan to integrate AI into battlefield management systems for troop coordination and enemy monitoring."

The system will initially be trained on secure Google GPU infrastructure abroad, before moving entirely to local servers to ensure sovereignty and protect against ongoing Russian strikes.

Gemma was selected after reviewing alternatives, including Meta's Llama, Mistral AI, and Chinese models such as DeepSeek and Qwen.

Data from over 90 government institutions will train the AI, providing richer context than global general-purpose models. Cybersecurity measures are being implemented to protect the system against prompt injection attacks and other threats. Kyivstar has installed 3500+ backup generators to maintain operations during Russian strikes.

Nvidia unveils open source tool for monitoring GPU health

Nvidia has introduced open source software designed to give data centre operators unprecedented visibility into GPU thermals and reliability as next-generation AI accelerators push cooling systems and infrastructure limits. The tool tracks power, temperature, airflow, utilisation, memory bandwidth, interconnect health, and other vital indicators across thousands of GPUs, helping enterprises manage escalating heat and reliability challenges.

Nvidia emphasised its commitment to transparency, stating: "It will include an open source client software agent — part of Nvidia's ongoing support of open, transparent software that helps customers get the most from their GPU-powered systems." The service is opt-in, customer-installed, and provides read-only telemetry, with the company stressing that its GPUs contain no hardware tracking features, kill switches, or backdoors.

The launch arrives as thermal stress becomes a defining constraint for AI infrastructure. Modern accelerators exceed 700W per GPU, with multi-GPU nodes reaching 6kW, creating dense heat zones and interconnect degradation risks. A Princeton University CITP study warns that such stress can shorten AI chip lifespan to one to two years, below the widely assumed one-to-three-year range.

The software offers a fleet-level dashboard for power draw, memory bandwidth, airflow patterns, errors, and configuration mismatches, enabling early detection of silent faults and bottlenecks.

