

Red Hat unveils developer preview of RHEL AI



Image Source: <https://developers.redhat.com>

Red Hat has unveiled a developer preview of Red Hat Enterprise Linux AI (RHEL AI), marking a significant step in the realm of artificial intelligence (AI) development. RHEL AI serves as a foundational model platform aimed at simplifying the creation, testing, and deployment of cutting-edge, open source Granite generative AI models, which are pivotal for powering enterprise applications.

Built upon the InstructLab open source project, RHEL AI integrates open source-licensed Granite large language models from IBM Research with InstructLab model alignment tools. These tools, based on the LAB (Large-scale Alignment for chatBots) methodology, are packaged into an optimised, bootable RHEL image, streamlining server deployments.

The primary objective behind RHEL AI and the InstructLab project is to empower domain experts to contribute directly to large language models (LLMs), thereby facilitating the development of AI-infused applications such as chatbots. RHEL AI offers a comprehensive solution by leveraging community-driven innovation, providing user-friendly software tools tailored for domain experts, and incorporating optimised AI hardware enablement.

Red Hat Enterprise Linux AI comprises four foundational components: Open Granite models, InstructLab model alignment, optimised bootable Red Hat Enterprise Linux, and enterprise support and indemnification.

To facilitate experimentation and eventual production deployment, Red Hat has outlined a phased approach, beginning with experimentation on local devices using the InstructLab CLI, followed by deployment on bare metal servers or cloud-based virtual machines, and culminating in production deployment at scale with OpenShift AI.

SensiML open sources its auto ML tools for TinyML

SensiML has open sourced a significant portion of its TinyML toolchain. CEO Chris Rogers is focused on how the company's open source auto ML tools can accelerate the adoption of ML in IoT devices. The company's flagship offering, the SensiML Analytics Toolkit, provides a comprehensive development platform for IoT devices, covering data collection, labelling, algorithm and firmware generation, and testing. By open sourcing Analytics Studio, it aims to foster community collaboration to enhance its features and facilitate the deployment of ML even in the smallest devices.

The open source initiative aligns with the broader trend in the tech industry towards greater transparency and community-driven development. SensiML's tools are designed to be hardware-agnostic, supporting various platforms, including Arm Cortex-M class microcontrollers, Intel x86-based CPUs, and QuickLogic's S3 platform.



Chris Rogers highlighted the challenges in the TinyML market, including resource constraints, the need for extensive data collection, and the skillset gap among developers. By leveraging open source models, SensiML hopes to address these challenges and accelerate the development of intelligent IoT devices. The company envisions a collaborative environment where developers can contribute to and benefit from advancements in TinyML technology.