Google announces LMEval to enhance large model evaluation

Google has launched LMEval, a powerful open source framework designed to make large model evaluation faster, easier, and more consistent across providers. This release aligns with its ongoing collaboration with Giskard, which uses LMEval to run the Phare benchmark—an independent test for assessing model safety and security.

With LLMs evolving rapidly, developers and researchers need reliable tools to assess model performance across tasks and providers. Until now, cross-model benchmarking has been complex and resource-heavy. LMEval changes that.

It is user-friendly, with example notebooks available in the GitHub repository. Running evaluations on different model versions—like Gemini—takes just a few lines of code. It comes with LMEvalboard, a dashboard for interactive model comparisons.

The companion dashboard to it enables users to explore and understand model performance in depth interactively. It allows for quick comparisons of overall model accuracy across benchmarks, helping users see how different models stack up. Users can also analyse individual models to uncover their strengths and weaknesses across various categories, gaining insight into specific performance trends. It makes it easy to perform head-to-head comparisons between models, highlighting areas of disagreement and revealing where one model may outperform another.

Databricks to acquire Neon, the serverless Postgres startup

Databricks has announced its intent to acquire Neon, a fast-growing serverless Postgres startup, in a move that could reshape the $100 billion-plus database industry amid a wave of AI-driven transformation.

As artificial intelligence redefines software development, Databricks aims to enhance its data intelligence capabilities by integrating Neon's cutting-edge, serverless Postgres platform that is purpose-built for modern, AI-driven applications.

Neon has rapidly emerged as a go-to platform for AI-native workloads. Internal telemetry from the company reveals a striking trend: over 80% of the databases provisioned on Neon were created automatically by AI agents and not humans. This signals a paradigm shift in how developers are building, deploying, and managing data systems.

"The era of AI-native, agent-driven applications is reshaping what a database must do," said Ali Ghodsi, co-founder and CEO of Databricks. "Neon proves it: four out of five databases on their platform are spun up by code, not humans. With Neon, we're giving developers a serverless Postgres solution that matches the speed, scalability, and openness that AI demands."

The integration of Neon's serverless Postgres architecture with the Databricks Data Intelligence Platform will enable enterprises to eliminate traditional compute-storage bottlenecks and streamline AI development workflows. The combined solution aims to simplify infrastructure, reduce costs, and accelerate innovation — all underpinned by Databricks' enterprise-grade security, governance, and scalability.

Nikita Shamgunov, CEO of Neon, expressed confidence in the partnership: "Four years ago, we set out to build the best cloud-native Postgres — serverless, scalable, and open to all. Joining Databricks supercharges our mission with the support of an AI leader. Together, we're entering a bold new phase."

Following the acquisition — pending standard regulatory approvals — Neon's team will join Databricks. The partnership is expected to help customers break down data silos, streamline AI architecture, and build more responsive, reliable, and secure AI agents.

ONLYOFFICE Docs 9.0 is here, packed with new features and enhancements

The highly anticipated ONLYOFFICE Docs 9.0 has been released, bringing over 20 new features and nearly 500 fixes designed to boost productivity. From a refreshed user interface to advanced AI tools and broader file format support, this update delivers a smoother, more powerful experience for all users.