

Power in Your Hands: Running Local Large Language Models for AI Brilliance

This step-by-step guide provides a good understanding of the Dalai Library and how to get started with the LLaMA and Alpaca language models on your own machine. Once you have done that, the applications are limited only by your imagination.



Figure 1: Image generated by Midjourney and further edited in Canva

Have you ever thought about running cutting-edge language models like ChatGPT on your own computer? You're not alone! In this adventure, we'll explore how to make it happen with the LLaMA and Alpaca models, using the Dalai Library. Get ready for a detailed, step-by-step, and entertaining journey into the world of local AI! At the end of it you will, hopefully, get a good understanding of how generative AI can reside in your machine.

The local AI dilemma: Our own private sanctuary

Let's face it, we're all excited about the possibilities

generative AI models bring. But with great power comes great responsibility (thanks, Uncle Ben). One of the biggest concerns we have is data privacy. Centralised models like OpenAI and Microsoft's offerings are fantastic, but do we really want to hand over our data on a silver platter?

Imagine if Batman had to share his Batcave location with everyone. Not cool, right? That's where running an AI model on your local machine comes into play. It's like having your very own Batcave (minus the cool gadgets and bat-themed vehicles, of course).