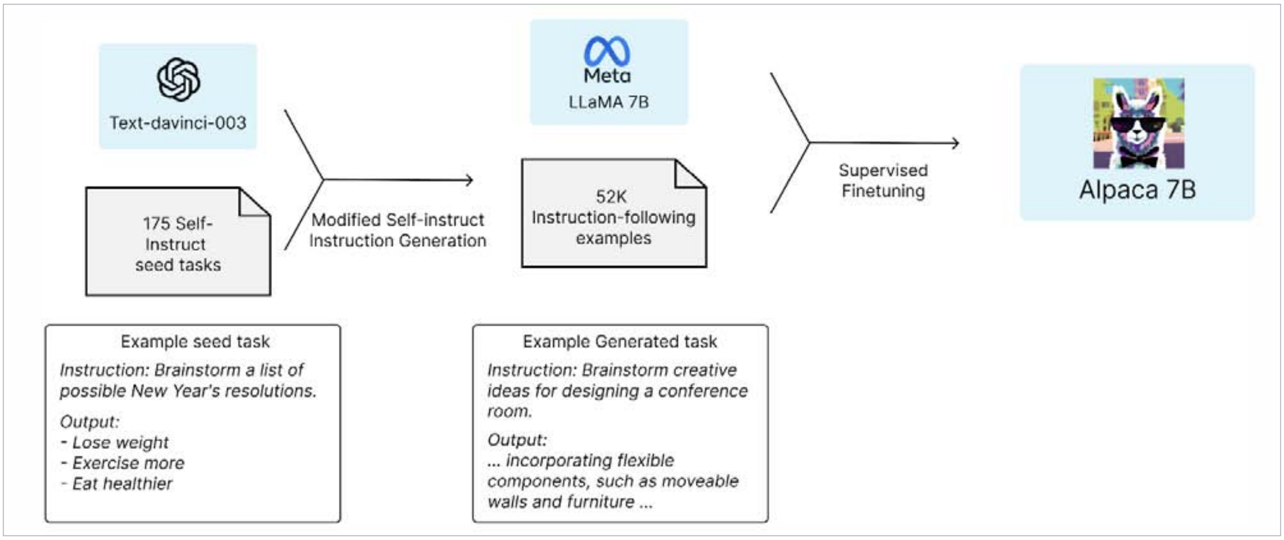

Figure 2: Paper published in the Center for Research on Foundation Models (CRFM) at Stanford

LLaMA: The compact powerhouse created by Meta AI

LLaMA is a foundational language model that has managed to achieve something incredible. Despite being 13x smaller than the colossal GPT-3, it outperforms the latter on most benchmarks! This compact powerhouse is capable of running on local machines – one daring individual even managed to get it working on a Raspberry!

Now, thanks to some unforeseen circumstances, LLaMA is available for non-commercial use. Developed by the talented team at Meta AI, LLaMA and its sibling Alpaca are making local AI usage more accessible than ever before. So, without further ado, let's get started on this great ride!

Installing and running LLaMA

Clone the repo and install the necessary prerequisites from *https://github.com/cocktailpeanut/dalai*.

To kick things off, run the command:

```
npx dalai llama install 7B
```

Before you proceed, though, be aware that LLaMA-7B needs around 31GB of storage. So make sure there's enough space on your computer for this small yet mighty guest. I struggled with this a lot!!

To run LLaMA, simply type:

```
npx dalai serves
```

..and you've done it! You now have a large language model running locally. Give yourself a pat on the back! Why? Because, I had to go through a lot of trouble to get this running like specific versions of Python and Node. You can find the details in the Readme file at *https://github.com/cocktailpeanut/dalai*.

However, when I ran this, there were a lot of meaningless machine letters that came in answer to a simple question "I feel like having some snacks because . ." Upon further research, it looks like this is a well-known issue. Someone in a discussion channel suggested: "Check Alpaca model," so I did that.

Alpaca: The instruction-following marvel

Alpaca is a fine-tuned version of LLaMA that's been designed to follow instructions, much like ChatGPT. The jaw-dropping part is that the entire fine-tuning process cost less than US$ 600! When you compare that to GPT-3's staggering US$ 5,000,000 price tag, it's an absolute steal!

So, how did they do it? Well, OpenAI's text-davinci-003 model lent a helping hand unwittingly by transforming 175 self-instruction tasks into a whopping 52,000 instruction-following examples for supervised fine-tuning. Talk about a clever workaround! The brains behind this amazing model are Rohan Taori, et al. It's so creative — they basically used da-vinci-03 as a teacher for LLaMA to produce AIpaca!! You can find the entire paper on this at *https://crfm.stanford.edu/2023/03/13/alpaca.html*.

Installing and running Alpaca

To install Alpaca, all you have to do is run:

```
npx dalai alpaca install 7B
```

Continuted to page...34