

alternative green open source repository. Anyone can download and use it, but paying customers receive version-specific patches. This is our business model. Customers who pay get the fix within 24 hours depending on the SLA and their agreement with us, while those who do not pay receive it whenever the open source community integrates the update.

Q. What role does the green open source repository play here? Does it only contain components provided by Persistent Systems, or do other external contributors also contribute to it?

A. Right now, we control the repository because it is tied to our service commitments and paying customers. However, we are expanding participation to build a developer community that allows anyone to contribute.

Developers using our SASVA platform will be able to create and publish fixes to this repository. We are actively contributing fixes to thousands of open source projects, and these fixes are automatically added to the green open source repository. Eventually, we aim to make this a primary repository, serving as an alternative for enterprises. If customers download software from our repository, they will always have the most up-to-date and secure version.

We will continue maintaining these repositories, ensuring if the community releases a fix, we provide it. If the community adds new features, we incorporate them. If there is a vulnerability or security risk, we fix it first, even before the community does.

Q. Red Hat, which is now part of IBM, used to offer a subscription model where subscribers would receive fixes for any vulnerabilities under a service level agreement (SLA). How does your approach compare to that, especially for open source software that does not have an SLA?

A. We are unique as we not only

fix open source vulnerabilities but also contribute those fixes back to the open source community and roll them out in an automated manner.

Many companies focus on identifying vulnerabilities and planning fixes, but in the open source world, vulnerabilities are discovered at a much faster rate than they are fixed, creating a significant security gap. We are working to bridge that gap by accelerating the remediation process. Our service benefits both customers and the global open source community. Our green open source repository is available to everyone, mostly for free. The only service we charge for is the SLA-based option—if a customer requires a fix within 24 hours, they pay for it. Otherwise, they can wait a few days for the fix to be publicly available.

Our vision is to use our home-built AI platform to maintain open source branches, provide rapid fixes, and keep critical products secure under a defined SLA. We act as an insurance policy for companies that build critical applications using open source, strengthening open source adoption across industries.

Q. Can you tell us more about how you utilise AI to detect vulnerabilities and provide fixes?

A. Detecting vulnerabilities is relatively easy as they are publicly disclosed—you can check the Common Vulnerabilities and Exposures (CVE) database, identify an issue, and inspect the affected code. The fixing part, however, is very complex where our SASVA platform comes into play. It links to any repository, and with a specific description of the issue, our AI-driven agent workflow analyses the entire repository to generate a fix. It produces a binary that meets Supply-chain Levels for Software Artifacts (SLSA) compliance and other open source testing standards. It delivers a final pull request to the open source community—within hours. We have already processed fixes for the top 100 open source repositories, covering issues from the last six months.

Now, we are scaling this capability to 10,000+ open source repositories worldwide. Our goal is to identify and fix vulnerabilities in real time.

Q. What makes SASVA different from any other platform?

A. Today, models have context-size constraints, but we have developed patented technology (over 35+ patents, with nearly 50 patents in progress) to bypass these context limitations. For example, Google Gemini has a 1-million-token context, while OpenAI models have 128KB (with newer models having context windows up to 1 million tokens). However, repositories can be as large as 15GB or 20GB. SASVA can scan and process entire repositories.

We have built SASVA entirely on an open source stack. Our training methodology ensures that our models remain adaptable to future AI advancements. SASVA understands the design patterns, frameworks, and code structures of open source projects. It runs on-premise, ensuring that data privacy is maintained. The platform is optimised for low-cost infrastructure. Our models can run on CPUs, GPUs, or even minimal hardware. This keeps hosting costs low, making it affordable and scalable for enterprises.

Q. Who is your target customer?

A. Our customers fall into several categories including large cloud providers, small and mid-sized enterprises, and major corporations developing their own applications.

One of the largest cloud hyperscalers relies on our solutions across all their cloud services. They require a strict SLA of 24 hours for security fixes. If a vulnerability similar to Log4j were to arise, they would receive a fix from us through their alternative green source, enabling them to provide timely solutions to their users. This is a B2B customer.

Another set of customers consists of enterprises that consume open