

Table 2: Key tools for performance evaluation and analytics of LLMs

Name	Core modules/ Capabilities	Key applications	URL
Comet Opik	Prompt/response logging, trace visualisation, experiment tracking, evaluations, OpenTelemetry support	LLM observability, evaluation and debugging	comet.com/opik
OpenLit	OpenTelemetry instrumentation, LLM spans, latency and token metrics, vendor-agnostic	LLM observability, tracing, prompt/response monitoring	openlit.io
OpenTelemetry	Traces, metrics, logs, exporters (Jaeger, Prometheus, console)	Unified observability backbone	opentelemetry.io
Helicone	Request proxying, latency analysis, prompt logs, error monitoring	LLM request monitoring and debugging	helicone.ai
Langfuse	Prompt versioning, trace trees, scoring, datasets, feedback loops	Prompt tracking, cost analysis, evaluations	langfuse.com
Promptfoo	Regression testing, prompt scoring, multi-model evals	Prompt testing and comparison	promptfoo.dev
Traceloop	OpenTelemetry-native tracing, framework integrations (LangChain, LlamaIndex)	LLM pipeline tracing	traceloop.com
Phoenix	Embedding drift, trace visualisation, evaluation dashboards	LLM quality and hallucination analysis	phoenix.arize.com/
Ragas	Faithfulness, relevance, context precision/recall	RAG and LLM output evaluation	github.com/explodinggradients/ragas
Evidently AI	Drift detection, data quality, evaluation reports	Data and model monitoring	evidentlyai.com
MLflow	Runs, metrics, artifacts, model registry	Experiment tracking (LLM-adaptable)	mlflow.org
DeepEval	Test cases, metrics, CI/CD-ready evaluation	Unit testing for LLMs	github.com/confident-ai/deepeval
Jaeger	End-to-end request tracing, span timelines	Trace visualisation	jaegertracing.io
Prometheus	Time-series metrics, alerting	Metrics monitoring	prometheus.io
Grafana	Traces, metrics, logs dashboards	Visualisation and dashboards	grafana.com
Ollama	Local model serving, OpenAI-compatible API	Offline/local LLM execution	ollama.com

India is a leader in the adoption and implementation of AI models.

LLM-as-a-judge: Observation, monitoring and analytics of LLM platforms

LLM-as-a-judge is a technology for evaluating the performance of AI LLMs in terms of effectiveness of prompts,

tokens used, cost factor, and many other parameters that directly affect the execution of the deployment.

The LLM-as-a-judge service provides observation and analytics with use cases centered on prompt lifecycle management—including assessing prompt effectiveness, ensuring traceability and versioning,

and analysing cost factors—alongside operational monitoring of latency and token accounting. Key advantages are robust governance and quality assurance, offering comprehensive response auditing, error diagnosis, and outcome visibility with capabilities for hallucination detection and regression testing. Furthermore, it enables