and deploy it on-premises, whether on physical machines, virtual machines, or using Kubernetes on containers. For users of AWS or Google Cloud, TiDB can be deployed on their underlying virtual machines (VMs) such as EC2 instances.

We offer managed services for cloud deployments. With TiDB Dedicated, users only need to point their applications to the service; we handle scalability, data maintenance, version upgrades, patches, backups, and recovery. Another option is TiDB Serverless, a pay-as-you-go model, which automatically scales up and down as per demand.

Q. How has TiDB adjusted its offerings for customers using AI and LLMs?

A. We are seeing a common pattern where most AI companies are leveraging vector databases. Although vector storage has existed for decades, its adoption has recently surged, especially with large language models (LLMs) and AI advancements.

Scalability is critical here, as vector storage grows alongside data and model accuracy requirements. TiDB ensures this scale without compromising performance, by integrating vector capabilities into the TiDB cluster. So TiDB can now support transactional workloads (OLTP), real-time analytical workloads, and vectorised data storage for AI models—all on a single cluster.

Currently, TiDB Vector Storage is available on TiDB Serverless and will be extended to on-premises and TiDB Cloud by the end of the year.

Q. What role does the community play in TiDB's offerings?

A. The open source community is instrumental in TiDB's growth. It boosts awareness, adoption, and innovation through shared experiences and feedback. We regularly conduct hackathons and receive creative contributions from the community. For example, someone once proposed implementing graph database capabilities on TiDB—a completely different use case that showcased the community's ingenuity.

In India, we see immense potential due to the availability of exceptional talent. We are actively investing in community-building efforts, hosting events, and planning initiatives like the TiDB User Day this year to further engage with the open source community.

Q. What governance model do you follow for community contributions? How do you manage external inputs?

A. Open source contributors can propose ideas, which are reviewed by our product team. This team evaluates the feature's value, prioritises it against our roadmap, and ensures alignment with TiDB's principles. Approved contributions go through rigorous testing across all supported versions to maintain quality. A centralised governance team acts as a gatekeeper to oversee this process and ensure compliance.

Q. Do you have a dedicated community management team, and is there one specific to India?

A. Yes, we have a dedicated global team for community management, including a local community manager in India. Given the talent pool here, we are expanding our efforts in India, hosting events and engaging with the open source community to foster collaboration and innovation.

Q. Are you involved in other open source projects?

A. One project is *OSSInsights.io*, an open source platform we developed initially to showcase TiDB's scalability by storing over 6 billion records in a single table. However, it quickly evolved as a tool for analysing trends in GitHub data, such as programming language popularity and database usage.

Q. Apart from hackathons and *OSSInsights.io*, what other ways do you engage with the community?

A. We have several engagement channels, including a Slack community for open source contributors. Recently, we introduced initiatives like offering free credits to contributors with innovative ideas for using TiDB. We also run contests to recognise the best TiDB use cases and stories, rewarding participants for their contributions. These activities are just the beginning; we aim to enhance our community engagement significantly in the coming years.

Q. How do you see distributed databases evolving and what can the community expect from TiDB in the next five years?

A. Distributed databases are poised for widespread adoption as businesses grow and demand scalable, efficient solutions. By 2030, the expected 10-20x growth in digital-native businesses, AI, and SaaS will make distributed databases a necessity rather than a choice.

The evolution of SQL-based distributed databases is a natural progression from traditional SQL and NoSQL databases. These modern databases combine SQL's familiarity with the scalability of NoSQL, ensuring they are future-proof for the growing demands of data-intensive industries.

For the open source community, this represents a tremendous opportunity for learning and innovation. TiDB is committed to being part of this journey, providing scalable, innovative solutions while remaining open to ideas and contributions from the community. Together, we can shape the future of distributed databases. END