

learning model using Scikit-learn is given below:

```
# Import necessary libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Load the Iris dataset
data = load_iris()
X, y = data.data, data.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Initialize and train the model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Make predictions
predictions = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, predictions)
print(f"Model Accuracy: {accuracy * 100:.2f}%")
```

Additional features

Pipeline creation: Scikit-learn allows chaining of preprocessing steps and modelling into a single pipeline using the Pipeline class. This ensures reproducibility and minimises code repetition.

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC

pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('svm', SVC(kernel='rbf'))
])
pipeline.fit(X_train, y_train)
pipeline_predictions = pipeline.predict(X_test)
```

Custom estimators: Users can create custom transformers and estimators to extend the library's functionality.

Use cases of Scikit-learn

Predictive analytics: Models can be built to predict outcomes based on historical data, such as customer churn, stock price forecasting, or disease outbreak prediction. Predictive analytics helps businesses make data-driven decisions and anticipate future trends.

Classification and regression: Helps solve problems like spam email detection, sentiment analysis, credit scoring, or price prediction. Algorithms like support vector machines (SVMs), logistic regression, and random forests are frequently used for these tasks.

```
from sklearn.linear_model import LogisticRegression

# Train a logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)

# Evaluate using classification report
from sklearn.metrics import classification_report
print(classification_report(y_test, predictions))
```

Clustering: Can group similar items, such as customer segmentation or document clustering, using K-means, DBSCAN, or agglomerative clustering. Clustering is widely used in marketing campaigns and recommendation systems to target specific user groups.

```
from sklearn.cluster import KMeans

# Perform K-Means clustering
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(X)
print(kmeans.labels_)
```

Dimensionality reduction: Helps reduce dataset size for visualisation or to improve model efficiency. PCA and t-SNE are particularly effective for visualising complex datasets in 2D or 3D. For example, reducing image feature dimensions can make computer vision tasks more efficient.

```
from sklearn.decomposition import PCA

# Apply PCA for dimensionality reduction
pca = PCA(n_components=2)
X_reduced = pca.fit_transform(X)
print(X_reduced)
```

Recommendation systems: Provides personalised recommendations using collaborative filtering or content-based methods. For instance, Scikit-learn can be used to build a recommendation system for e-commerce platforms, movie streaming services, or online learning platforms.

Anomaly detection: Identifies outliers or rare events using algorithms like Isolation Forest or One-Class SVM. This is especially useful in fraud detection, network security, and industrial equipment monitoring.