stopwords

A number of stopword lists are supported such as 'catalan', 'romanian', 'german', and 'SMART' (for English). You can list the stopwords defined for the list using the 'kind' argument to the *stopwords()* function, as shown below:

```
> stopwords(kind = "en")
  [1] "i"       "me"         "my"         "myself"     "we"
  [6] "our"    "ours"        "ourselves"  "you"        "your"
 [11] "yours"  "yourself"   "yourselves" "he"         "him"
 [16] "his"    "himself"    "she"        "her"        "hers"
 [21] "herself"    "it"     "its"        "itself"     "they"
 [26] "them"    "their"    "theirs"     "themselves" "what"
```

stripWhitespace

Any extra whitespace can be removed from the input text using the *stripWhitespace()* function. For example:

```
> l <- removeWords(lines[24], stopwords("en"))

> l
[1] "instance,   early years  automobile development  group
capital monopolists owned  rights  "

> stripWhitespace(l)
[1] "instance, early years automobile development group

capital monopolists owned rights "
```

tokenizer

The text document can be tokenized using a tokenization algorithm. The *Boost_tokenizer(), MC_tokenizer()* and *scan_tokenizer()* functions are available. An example of the *Boost_tokenizer()* function is given below:

```
> getTokenizers()
[1] "Boost_tokenizer" "MC_tokenizer"    "scan_tokenizer"

> Boost_tokenizer(lines[24])
 [1] "instance,"    "in"     "the"        "early"
"years"
 [6] "of"    "automobile"  "development" "a"
"group"
[11] "of"     "capital"     "monopolists" "owned"
"the"

[16] "rights"      "to"          "a"
```

weightBin

A binary weight on a term-document matrix can be computed using the *weightBin()* function, as demonstrated below:

```
> tdm <- TermDocumentMatrix(lines)

> weightBin(tdm)
<<TermDocumentMatrix (terms: 4167, documents: 1465)>>
Non-/sparse entries: 9623/6095032
Sparsity           : 100%
Maximal term length: 139
Weighting          : binary (bin)
```

inspect

You can view a detailed information on the corpus, term-document matrix or input text using the *inspect()* function, as illustrated below:

```
> inspect(tdm)
<<TermDocumentMatrix (terms: 4167, documents: 1465)>>
Non-/sparse entries: 9623/6095032
Sparsity           : 100%
Maximal term length: 139
Weighting          : term frequency (tf)
Sample             :
             Docs
Terms         1183 32 340 37 529 759 82 828 83 971
  and            0  2   0  3   1   0  0   1  0   0
  for            1  0   0  0   2   1  1   0  1   1
  free           0  0   0  0   0   0  0   0  0   0
  open           1  0   0  0   0   0  0   0  0   0
  open-source    0  0   2  0   0   0  0   0  0   0
  retrieved      0  0   0  0   0   0  0   0  0   0
  software       0  0   1  0   0   0  0   0  0   0
  source         1  0   0  0   1   0  0   0  0   0
  that           0  0   0  1   0   0  0   0  0   0
```

writeCorpus

The *writeCorpus()* function can output a text format of the corpus to multiple files, as illustrated below:

```
> writeCorpus(feed, "/tmp/foo")

$  ls /tmp/foo/
1000.txt  1061.txt  1121.txt  1182.txt  1242.txt
...
```

You are encouraged to read the manual page of the tm package to learn more about its functions, arguments, options and usage. END

By: Shakthi Kannan

The author is a free software enthusiast.