

AI Gatekeeper offers real-time security for autonomous agents



Operant AI's CEO Vrajesh Bhavsar (extreme right) with his colleagues (from left) Ashley Roof and Priyanka Tembey

Silicon Valley-based Operant AI has introduced AI Gatekeeper, a real-time security framework designed to defend live AI applications, autonomous agents, and complex agentic AI workflows — across Kubernetes, private clouds, hybrid setups, and edge environments. With AI-native systems evolving beyond traditional models, especially in growth markets like India, enterprises are rapidly deploying autonomous AI agents.

According to Deloitte, over 80% of Indian companies are actively exploring agent-based AI solutions, while 50% are scaling multi-agent workflows — all with limited human oversight. This shift introduces new security challenges that existing cloud and AI defences aren't built to handle. Operant's AI Gatekeeper steps in as an open, runtime solution built for the AI-native era, extending beyond traditional perimeter security.

As enterprises increasingly rely on third-party vendors and hyperscale platforms (AWS, Azure, GCP, Databricks, Snowflake, Salesforce), Operant AI's solution addresses risks like model poisoning, data leakage, and unauthorised agent actions — issues amplified by the rise of autonomous workflows.

AI Gatekeeper's release closely follows Operant's inclusion in Gartner's AI TRiSM (AI Trust, Risk, and Security Management) Market Guide — positioning it as a key player in the security-first AI development space.

With AI workflows now expanding to wherever the data resides, tools like AI Gatekeeper are set to become foundational for securing the future of intelligent applications. "Our AI security problem today isn't what it was two years ago," said Vrajesh Bhavsar, CEO at Operant AI. "Agentic AI creates a dynamic attack surface that needs real-time, distributed protection."

Open source 3D-printed robot for automating chemical experiments

A team of scientists from Japan's Hokkaido University, led by Professor Keisuke Takahashi, has developed FLUID (Flowing Liquid Utilizing Interactive

Mozilla's new open source tools help developers build ethical AI datasets

Mozilla has unveiled open source tools to help developers build ethical AI datasets and avoid training models on copyrighted material.

The reliance of many popular large language models (LLMs) on vast datasets scraped from the internet, often encompassing copyrighted works used without permission, presents a significant ethical and legal challenge.

A growing contingent within the developer community believes creating high-quality, ethically sound alternatives is not only possible but necessary. This launch directly supports that movement.

The new toolkits — products of a year-long collaboration with EleutherAI — provide developers with practical workflows, code, and demonstrations hosted on the Mozilla.ai Blueprints platform (a space dedicated to helping developers prototype AI applications using open source components).



Ayah Bdeir, senior advisor for AI strategy at the Mozilla Foundation, said: "Just like open source software in its early days, today's open data ecosystem depends on community contributions and shared values. These toolkits are part of an effort to create common resources, make dataset creation easier, and promote the infrastructure needed for ethical AI development. Our partnership with EleutherAI is grounded in our shared commitment to advance this mission."