

Cloudflare open sources Pipit to democratise large AI models

Cloudflare has open sourced Project Pipit, releasing a lossless compression tool that could fundamentally reshape how large language models are distributed and deployed. Unlike traditional approaches such as quantisation and pruning—which compromise accuracy to reduce size—Pipit compresses models without altering a single numerical value.

Outputs, probability distributions, and benchmark scores remain byte-for-byte identical after decompression, eliminating the long-standing trade-off between efficiency and fidelity. Developed under Dr Adaosa Okafor, director of Machine Learning at Cloudflare, the tool uses a proprietary entropy-coding algorithm to achieve fully reversible compression. Benchmarks show compression ratios of approximately 5.2× on dense Llama-3 class models and 3.8× on Mixture-of-Experts architectures.

Integrated with Cloudflare Workers AI, Pipit enables models to be stored in compressed form, streamed from edge locations, and decompressed with near-zero latency. For models exceeding 70 billion parameters, reduced network transfer times outweigh the minimal runtime overhead.

The tool is designed for immediate adoption, supporting PyTorch and SafeTensor formats without requiring changes to existing pipelines. Beyond technical gains, the open source release signals a strategic shift. By reducing storage and bandwidth costs by up to five times, Pipit weakens dependence on centralised GPU cloud providers.

If widely adopted, Pipit could emerge as a standard compression layer for AI models, accelerating decentralised AI and redefining global model distribution.

Concerns deepen because MizarVision reportedly has 5.5% Chinese government ownership and has also published imagery tied to Israeli positions, Diego Garcia deployments, Australian naval movements, and even TSMC’s semiconductor plant construction.

NVIDIA launches Ising, the world’s first quantum AI model family

NVIDIA has unveiled Ising, the world’s first family of open source quantum AI models, positioning open source AI as the control layer for fault-tolerant quantum computing infrastructure.

The launch directly addresses the two biggest engineering roadblocks in scaling quantum systems—processor calibration and quantum error correction—with performance claims of up to 2.5x faster decoding and 3x higher accuracy than pyMatching, the current open source benchmark.

The Ising family includes Ising Calibration, a vision-language model that automates processor tuning, and Ising Decoding, two 3D convolutional neural network variants optimised for either speed or accuracy in real-time quantum error correction. NVIDIA said the calibration workflow can now shrink from days to hours through AI agents performing continuous automated tuning.

“AI is essential to making quantum computing practical. With Ising, AI becomes the control plane — the operating system of quantum machines — transforming fragile qubits to scalable and reliable quantum-GPU systems,” said Jensen Huang, founder and CEO, NVIDIA.

The open source release arrives with a quantum workflow cookbook, training datasets, NVIDIA NIM microservices, local deployment support, and hardware-specific fine-tuning tools, enabling researchers to adapt models while preserving proprietary data.

Adoption at launch already spans Academia Sinica, Fermi National Accelerator Laboratory, Harvard, IonQ, IQM Quantum Computers, Sandia National Laboratories, UC San Diego, and Yonsei University, among others.

Integrated with CUDA-Q and NVQLink, and distributed through GitHub, Hugging Face, and build.nvidia.com, Ising marks a major open source infrastructure play as the quantum computing market races towards \$11 billion by 2030.

WSO2’s API Platform will help enterprises enable agentic AI

WSO2 has announced the general availability of its API Platform, an open source solution built on an Apache 2.0 foundation, designed to unify the management of traditional APIs and AI assets from a single control plane. The platform is positioned to eliminate vendor lock-in through a modular architecture while supporting self-managed SaaS and hybrid deployments, alongside multi-gateway federation across ecosystems such as Kong, Amazon API Gateway, and Azure API Management.

At its core, the platform enables agentic AI consumption by making APIs accessible to autonomous agents via the Model Context Protocol (MCP). It delivers end-to-end lifecycle management, covering publishing, versioning, and discovery of both APIs and AI tools, while embedding governance for cost, risk, and compliance.

A key innovation is the AI Gateway, which converts APIs into agent-ready tools within minutes and governs both internal and third-party MCP services. Features such as semantic caching, adaptive routing, token-based rate limiting, and model routing help control LLM access and prevent data leakage.

The AI Workspace introduces centralised governance with over 30 built-in guardrails through a Policy Hub, while allowing custom guardrails in Go and