


Why Decentralised Identity is the Security Bedrock for Agentic AI

As AI agents begin to act on behalf of executives (signing contracts, moving funds), how do we verify their identity? In the framework proposed here, open source decentralised identifiers act as the ‘passport’ for AI agents, ensuring every autonomous action is cryptographically signed and authorised.



In the rapid transition towards an AI-driven enterprise, leadership focus has largely remained on model performance and accuracy. However, in 2026, a more systemic risk has emerged: the ‘accountability gap’. As organisations move from ‘assistive AI’ (chatbots) to ‘agentic AI’ (autonomous actors), these digital entities are being empowered to negotiate contracts, access sensitive repositories, and execute financial transactions.

Traditional identity and access management (IAM) systems, designed for human users and static service

accounts, are proving insufficient for a workforce of ephemeral, high-velocity AI agents. To secure the modern enterprise, architects must look towards decentralised identity (DID). By providing every AI agent with a cryptographically verifiable ‘passport’, we can ensure that every autonomous action is authorised, auditable, and revocable in real-time.

The ‘who’ problem in the age of ‘what’

In the current AI gold rush, enterprises have focused on what AI can do—

generate code, summarise reports, or automate customer service. But as we move into the era of agentic AI, where autonomous models are empowered to move funds, sign contracts, and access sensitive PII, a far more critical question emerges: ‘who’ is performing this action, and ‘who’ authorised it?

For technology organisations, the risk is no longer just a ‘hallucination’; it’s an ‘unauthorised execution’. Traditional identity systems were built for humans and static service accounts, not for autonomous agents that mutate and scale. To scale AI