




Figure 1: Lifecycle of an autonomous request: The 'trust-before-task' sequence

Do	Don't
Do prioritise W3C standards for DIDs to ensure future interoperability between different AI agents.	Don't use static API keys for AI agents; they are easily leaked and lack granular permissioning.
Do start with 'read-only' agents to prove the identity framework before moving to 'write/execute' agents.	Don't build on closed source identity silos that charge 'per-agent' fees; the costs will scale exponentially.
Do use Zero-Knowledge Proofs (ZKPs) to allow agents to prove they have authorisation without revealing sensitive backend data.	Don't ignore the 'liveness' problem—ensure your architecture can rotate or revoke an AI's identity instantly.

The executive's guide: Dos and don'ts

inherently trusted based on its network location, but only through its proven, decentralised identity.

A robust architecture for decentralised intelligence separates the 'brain' (the LLM) from the 'passport' (the DID).

The identity layer: A distributed ledger (e.g., Hyperledger Fabric) acts

as the 'root of trust', hosting the public keys (DID documents).

The agent wallet: Each agent instance has a sidecar 'wallet' (Aries) that handles cryptographic signing.

The policy engine: Before an agent executes a tool call, an open source governance engine verifies the agent's credentials against corporate policy.

Implementation guide: Dos and Don'ts

To implement this architecture effectively, leadership must prioritise the issuance of short-lived, task-specific credentials rather than persistent administrative rights, as AI agents are ephemeral, and their access should expire the moment a mission is completed. It is vital to embed model explainability alongside identity so that every cryptographically signed action is accompanied by a 'reasoning trace' for future audits. Yet one must never rely on static API keys or hard-coded secrets which represent the single greatest vulnerability in automated workflows. Furthermore, while the goal is full autonomy, architects should avoid removing the human-in-the-loop for high-value transactions; instead, use multi-signature requirements where an