

Microsoft to grant Azure credits to open source projects



Microsoft has announced a programme that will grant Azure credits to open source projects for a year. According to company officials, any project in any technology with an Open Source Initiative (OSI) approved licence may apply.

Developers will be able to use these credits for testing, storage, builds, and other development work.

“Open source software is an integral part of development at Microsoft, aligned with our goal to empower all developers to be successful in building any application, using any language, on any platform. We are committed to building open, flexible technology and working with the open source community to grow together as an industry. We are giving back to the open source ecosystem we participate in and depend on,” wrote Carol Smith, senior program manager, Microsoft Open Source Programs Office in a blog post.

Microsoft has already extended these credit grants to some projects in the open source ecosystem including FreeBSD, Alma Linux, Haskell, Snakemake workflow management and Promitor.

Those interested can apply through the Azure Credits website. Each grant will be for one year. Applicants will need to re-apply each year if they would like credits again.

Red Hat announces Red Hat Ansible Automation Platform 2

Red Hat, Inc. has announced Red Hat Ansible Automation Platform 2 as the company’s new standard for hybrid cloud automation. The latest version of the platform adds self-contained automation capabilities while shifting automation more deeply into the application development life cycle.

Google Cloud announces Backup for Google Kubernetes Engine

Google has announced the preview of Backup for GKE, a cloud-native way to protect, manage, and restore containerised applications and data running on Kubernetes.

Using the new service, developers can create plans to schedule periodic backups of both application data and GKE cluster state data, restoring each backup to a cluster in the same or in a different region. It also allows custom backups to ensure application consistency for the most demanding, tier-one database workloads.

Organisations are choosing to build on Google Kubernetes Engine (GKE), driven by benefits like higher developer productivity and lower infrastructure costs. And one of the fastest growing GKE architectures is the deployment of stateful workloads like relational databases inside GKE containers.

Chris Schilling and Manu Batra, product managers at Google Cloud, explained: “Prior to Backup for GKE, many GKE customers backed up their stateful application data separately from GKE cluster state data. Application data could be protected via a storage-based backup, while cluster state data might be captured occasionally using custom scripts and stored in a separate customer bucket. Customers with ongoing backup requirements relied on homegrown solutions to perform regular backups and to demonstrate compliance. In the event of a restore, customers had to perform more complex orchestration. Storage management tasks, like creating a clone for testing purposes, or migrating data from one cluster to another, meant additional operational overhead.”