

SUSE announces India expansion plans



SUSE, a global leader in enterprise-grade open source solutions, plans to work with large institutions in India to create centres of excellence in platform engineering to address the gap in talent and expertise around modern data centre technologies like Kubernetes and cloud-native applications.

The centres of excellence will provide electives and project opportunities for students in areas like data centre engineering and programmable data centres to build a pipeline of skilled talent familiar with SUSE's open source solutions.

"Our goal is to nurture local talent through internships and direct collaboration with our customers in cutting-edge technologies to support India's growing demand for cloud-native and open source solutions," said Rajarshi Bhattacharyya, country manager for India, SUSE.

As part of this expansion, SUSE will integrate the recently acquired StackState, a full-stack observability platform, into its Rancher Prime platform, to provide Indian enterprises with enhanced visibility into their containerised environments, enabling faster detection of issues and potential security threats.

Vishal Ghariwala, senior director and CTO for Asia Pacific at SUSE, declared the company's intentions to open source the StackState platform soon. "SUSE's open source ethos resonates well in India, where there is a growing culture of transparency and collaboration," he said.

Interestingly, the National Payments Corporation of India (NPCI), responsible for processing over 18 billion transactions monthly, has adopted SUSE's Rancher Prime and Liberty Linux solutions to manage its vast, containerised environments and enhance security.

Shiva Kumar RV, deputy CTO for infrastructure, National Payments Corporation of India (NPCI) told *OSFY* that "SUSE's Rancher Prime platform provides NPCI with a centralised Kubernetes management solution, allowing the latter to scale operations efficiently while maintaining stringent security protocols."

RHEL's next-gen AI platform supports GPUs

Red Hat has officially released Red Hat Enterprise Linux AI (RHEL AI) 1.2, advancing its foundation model platform for generative AI (GenAI). RHEL AI focuses on developing, testing, and running large language models (LLMs) for enterprise applications using Granite LLMs and InstructLab tools, distinct from the core RHEL operating system.

RHEL AI introduces AMD Instinct GPU support for the first time. With the integration of the ROCm open source compute stack, organisations can leverage MI300X GPUs for both training and inference and MI210 GPUs for inference-only workloads. This addition enhances flexibility by complementing existing NVIDIA GPU and CPU-based operations.

RHEL AI is now accessible on Microsoft Azure and Google Cloud Platform, enabling enterprises to deploy AI-based GPU instances effortlessly. This cloud support allows businesses to run AI workloads seamlessly across multi-cloud environments. Users can now deploy it on Lenovo ThinkSystem SR675 V3 servers. Preloaded deployment options simplify the setup process, ensuring better compatibility with hardware accelerators and enhancing overall efficiency. A new periodic checkpointing feature allows users to save long training sessions at specific intervals. This capability ensures that models can be fine-tuned from the latest checkpoint, improving efficiency and conserving computational resources.

The `ilab` CLI now includes automatic detection of GPUs and other hardware accelerators. This feature minimises manual configuration efforts, optimising performance by aligning with the available hardware. It introduces fully sharded data parallel (FSDP) via PyTorch, enabling distributed model training by sharding parameters, gradients, and optimiser states. This feature significantly reduces training time for complex models, enhancing overall productivity.