

ASF elevates Apache DevLake and Grails to top-level status

The Apache Software Foundation (ASF) has announced that Apache DevLake and Apache Grails have graduated to Top-Level Projects (TLPs), signalling maturity, community growth, and operational independence.

Apache DevLake, launched in December 2021, is an open source development data platform that integrates, analyses, and visualises fragmented DevOps data. By unifying metrics from various tools, DevLake helps software teams generate actionable insights to enhance workflow efficiency and decision-making.

Apache Grails, a web application framework built on Apache Groovy and Spring Boot for the JVM, celebrates its 20th anniversary with renewed momentum after revitalisation during its incubation period. James Fredley, Apache Grails PMC chair, said, “Becoming an ASF Top-Level Project signals the beginning of a new chapter for Apache Grails. With ASF’s support and a thriving contributor community, we look forward to broadening adoption and advancing the project’s capabilities.”

ASF’s governance framework emphasises community-led development and sustainability. Projects begin as ‘podlings’ in the Apache Incubator, where they receive mentorship, technical guidance, and community-building support. Elevation to TLP status indicates readiness for long-term success under The Apache Way.



Positioned as a direct rival to DeepSeek and OpenAI, Ling-1T underscores China’s growing emphasis on transparent, community-accessible AI ecosystems. The model, which Ant describes as a general-purpose LLM, delivers exceptional performance in code generation, software development, competition-level mathematics, and logical reasoning.



According to the Hangzhou-based firm, Ling-1T demonstrated “superior complex reasoning ability and overall advantage” compared with both open source and closed source models. It surpassed DeepSeek-V3.1-Terminus, Moonshot AI’s Kimi-K2-0905, and OpenAI’s GPT-5-main across benchmark platforms, including LiveCodeBench and the American Invitational Mathematics Examination (AIME).

In the AIME benchmark, Ling-1T achieved an accuracy of 70.42 per cent at an average cost of over 4000 output tokens per problem—performance that places it on par with Google’s Gemini-2.5-Pro while outperforming DeepSeek, OpenAI, and Moonshot models.

This is Ant’s second trillion-parameter model, following Ring-1T-preview, which it earlier claimed as the world’s first open source trillion-parameter thinking model. Ant first entered the AI model race in 2023 with its self-developed financial LLM.

By releasing Ling-1T as open source, Ant signals a strategic shift towards openness in AI development, strengthening China’s role in the global AI landscape and directly challenging the dominance of proprietary western AI systems.

Open source Nezha tool weaponised by hackers to deploy Gh0st RAT

Threat actors with suspected ties to China have weaponised the legitimate open source monitoring tool Nezha to deliver Gh0st RAT, turning a benign OSS project into a potent attack platform. The campaign, observed by cybersecurity company Huntress in August 2025, reportedly compromised over 100 machines, with the highest concentration of infections in Taiwan, Japan, South Korea, and Hong Kong, and additional victims across Singapore, Malaysia, India, the UK, the US, Colombia, Laos, Thailand, Australia, Indonesia, France, Canada, Argentina, Sri Lanka, the Philippines, Ireland, Kenya, and Macao.

Attackers initially exploited a publicly exposed, vulnerable phpMyAdmin panel to gain access. They used an unusual log poisoning technique, writing a one-liner PHP web shell into server logs and naming the file with a .php extension for direct execution. A Huntress spokesperson explained: “They then issued a query containing their one-liner PHP web shell, causing it to be recorded in the log file. Crucially, they set the log file’s name with a .php extension, allowing it to be executed directly by sending POST requests to the server.”