

Accubits' GenZ 70B LLM achieves top global ranking

Accubits Technologies, headquartered in Thiruvananthapuram, India, has taken a giant stride in the world of artificial intelligence (AI) by open sourcing its GenZ 70B large language model (LLM). This collaborative effort with Bud Ecosystem, another Accubits company, has resulted in GenZ 70B achieving the remarkable status of being the top-ranking instruction-tuned LLM on HuggingFace's



global leaderboard. Additionally, it stands at a commendable sixth place among open LLMs in all categories.

This 70-billion-parameter fine-tuned LLM represents a significant milestone in the AI landscape. Notably, it is positioned

as the largest LLM model in the OpenAI GPT series, making it a key player in the evolving field of generative AI.

Aharsh MS, chief marketing officer at Accubits Technologies, commented on the achievement, emphasising the local significance of GenZ 70B's global recognition. He stated, "An LLM model from India ranking top on a global scale is significant and can be an inspiration for the local developer community."

GenZ 70B is not just a run-of-the-mill language model; it boasts an autoregressive design with an optimised transformer architecture. To achieve its superior performance, it underwent fine-tuning using the supervised fine-tuning (SFT) method, incorporating data sets from OpenAssistant's instruction fine-tuning and Thought Source's Chain of Thought (CoT) approach. This rigorous fine-tuning has endowed GenZ 70B with enhanced skills and capabilities beyond those of a pre-trained model.

In terms of performance, GenZ 70B scored 7.34 on the MT-Bench benchmark, coming close to the performance of GPT 3.5, despite being 2.5 times smaller. This efficiency is particularly advantageous for real-world commercial applications, where smaller models are often more cost-effective and easier to work with.

Ruby on Rails creator criticises 'open source hooliganism'

Ruby on Rails creator, David Heinemeier Hansson, has expressed his concerns regarding what he calls 'open source hooliganism'. This critique arises from a recent incident involving the TypeScript community's reaction to a decision made by the Turbo team.

Hansson, well-respected in the open source realm, acknowledged the passion displayed by developers for their preferred programming languages and tools.



However, he cautioned against crossing the line between passionate advocacy and dogmatic crusading.

The incident in question centres on Turbo's choice to remove TypeScript from their project. While legitimate

concerns were raised, most of the debate was dominated by what Hansson termed 'open source hooligans' — individuals with no direct stake in Turbo but fervently defending TypeScript and attacking dissenting voices.

Hansson detailed this situation, citing numerous pull requests requesting TypeScript's reinstatement into Turbo and even a demand for his removal from the

LinkedIn's 'iris-message-processor' boosts escalation management

LinkedIn has unveiled its 'iris-message-processor', an open source project set to improve its Iris escalation management system's performance and reliability. This development offers faster processing speeds—around 4.6 times faster under average loads and an impressive 86.6 times faster under high loads compared to its predecessor.

The 'iris-message-processor' replaces the previous iris-sender subprocess with a fully distributed system, enabling more concurrent processing and direct message sending to vendors. It reduces reliance on the database, no longer using it as a message queue.

Notable improvements include the system's ability to handle large escalations, processing 6000 escalations in under 10 seconds, a significant leap from the previous 30-minute processing time. Even with 50% of nodes dropping simultaneously, it rebalances the cluster in under 30 seconds and processes active escalations in less than three seconds.

This performance boost is achieved by splitting Iris escalations into buckets, assigned dynamically to different nodes within the iris-message-processor cluster. The system's code is written in Go, designed to be efficient and scalable while reducing strain on the database.

LinkedIn has successfully used the iris-message-processor for a year with no outages, consistently surpassing SLOs (service level objectives). It offers a drop-in replacement for the iris-sender, retaining the existing Iris API while delivering substantial performance improvements.

The company has made the iris-message-processor an open source project, making it accessible to the wider community. It operates seamlessly in external environments, potentially replacing commercial incident management systems.