

GNU Linux-Libre 6.11 kernel released

The GNU Linux-libre project has announced the release and general availability of the GNU Linux-libre 6.11 kernel, aimed at users seeking 100% freedom for their GNU/Linux systems.

Built on the recently released Linux 6.11 kernel, the GNU Linux-libre 6.11 kernel cleans up drivers such as AMDGPU isp, tn40, rtl8192du, cs40l50, rt1320, and pcie-rcar-gen4, along with new AArch64 (ARM64) devicetree files. It also revises the cleaning of drivers like AMDGPU, Adreno, vgx61, atomisp, btnxpuart, prueth, and tas2781, adds a nowait-nowarn version of firmware loading/rejecting primitives, and introduces firmware-loading primitive detection in Rust.

This release includes all the features, enhancements, and updated drivers present in the Linux 6.11 kernel series, but with non-free components removed. Unlike the upstream kernel, the GNU Linux-libre version ensures that no proprietary code is included.

The GNU Linux-libre kernel is ideal for users who want a 100% free computer without any proprietary software. If you're one of those users, you can download the GNU Linux-libre 6.11 kernel now as compressed tarballs from the official website.

Packages for Debian-based (DEB) and Red Hat-based (RPM) distributions are also available via the Freesh project (DEB) and RPM Freedom



(RPM) websites. The GNU Linux-libre kernel can be installed on virtually any GNU/Linux distribution, either alongside or as a replacement for your distribution's default kernel.

API service powered by its SN40L AI chip, aimed at accelerating inference tasks for developers. The service supports Meta's Llama 3.1 models, including the largest 405-billion-parameter variant.

The introduction of this service highlights the increasing importance of inference in generative AI. The company pointed to growing demand for alternatives to NVIDIA, especially from enterprises seeking faster inference speeds. Inference, which involves running AI models after training, presents an opportunity for companies like SambaNova.

Recently, smaller AI chip makers, including SambaNova, have launched products focused on the inferencing market, offering claimed advantages over GPUs from NVIDIA, AMD, and others.

The SN40L chip features a three-tier memory system that combines SRAM, HBM, and DDR5, allowing it to handle both training and inference workloads efficiently. The company claims that SambaNova's architecture is more efficient than competitors that rely heavily on SRAM.

SambaNova's cloud service runs the 405-billion-parameter Llama model at over 100 tokens per second, surpassing Groq and Cerebras in speed and efficiency. The global AI chip market, forecasted to grow significantly, offers opportunities for new players as demand for NVIDIA's H100 creates supply shortages.

The Linux Foundation Decentralised Trust: Collaborating for a digital-first economy

The newly formed Linux Foundation Decentralised Trust seeks to encourage collaboration and innovation within the Web3 ecosystem, encompassing blockchain, ledgers, identity, interoperability, and cryptographic technologies.

With more than 100 founding members, the LF Decentralised Trust positions itself as a neutral platform for the joint development of technologies that support the shift towards a digital-first economy. The organisation builds on over eight years of work from the Linux Foundation, including prominent global open source projects and standards related to blockchain and digital identity.

Daniela Barbosa, general manager of decentralised technologies at the Linux Foundation and executive director of LF Decentralised Trust, said: "Around the world, decentralised

technologies are modernising critical systems and core infrastructure. LF Decentralised Trust is the new home for collaboratively developing the ecosystems that will make these trusted systems and applications."

The new organisation integrates all projects, labs, and communities from the Hyperledger ecosystem, including 200 local, regional, and industry groups that bring together tens of thousands of participants worldwide. In addition, TrustOverIP (ToIP) is now part of the fold, contributing its working groups, specifications, models, tools, and best practices for the governance of digital trust ecosystems as a project under the organisation's umbrella.

Michael Klein, blockchain technology lead at Accenture and chair of the LF Decentralised Trust Governing Board, added: "As we look to the future of digital ownership, open standards and open source initiatives will drive cross-industry collaboration and serve as the foundation for the next generation of the Internet."

LF Decentralised Trust launches with a total of 17 projects, including the entire Hedera Hashgraph codebase, which has been contributed as the LF Decentralised

