SUSE incorporates observability into SUSE Rancher Prime

SUSE has announced the general availability of SUSE Rancher Prime 3.1, which introduces new observability capabilities. This addition goes beyond typical infrastructure management, offering full-stack visibility that enables organisations to proactively identify and resolve issues, minimising downtime and preventing disruptions to operations and revenue.

With real-time, context-rich insights, Rancher Prime streamlines troubleshooting, allowing IT teams to focus on innovation rather than crisis management. It helps maintain optimal performance from edge to cloud, ensuring smoother, more reliable user experiences. Leveraging Rancher Prime's observability, businesses can optimise infrastructure, reduce operational risks, and improve agility.

These enhanced capabilities stem from SUSE's acquisition of observability vendor StackState in June and meet the growing needs of organisations as they manage increasingly complex, cloud-native environments. "As the threats of service interruptions continue, the risk of widespread outages is higher than ever and our customers need a unified view of their entire mission-critical infrastructure and applications," said Peter Smails, SVP, general manager, enterprise container management at SUSE. "Rancher Prime Observability will enable customers to ensure reliability, accelerate troubleshooting, and maximise performance while keeping their systems running smoothly."

With the addition of observability, Rancher Prime now offers a comprehensive set of capabilities that provide full visibility into everything managed by the platform. Available in both SaaS-hosted and on-premises deployment options, it enhances operational insights and efficiency for companies using Rancher Prime at scale.

Red Hat Enterprise Linux AI generally available for hybrid cloud environments

Red Hat, Inc., has announced the general availability of Red Hat Enterprise Linux (RHEL) AI for hybrid cloud environments. RHEL AI serves as Red Hat's platform for foundation models, allowing users to efficiently develop, test, and deploy generative AI (gen AI) models to enhance enterprise applications. This platform integrates the open source Granite family of large language models (LLMs) and the InstructLab model alignment tools, utilising the Large-scale Alignment for chatBots (LAB) methodology. It is delivered as an optimised, bootable RHEL image, designed for individual server deployments across hybrid cloud systems.

Image Source: https:www.redhat.com

"RHEL AI provides the ability for domain experts, not just data scientists, to contribute to a built-for-purpose gen AI model across the hybrid cloud, while also enabling IT organisations to scale these models for production through Red Hat OpenShift AI," said Joe Fernandes, vice president and general manager, Foundation Model Platforms, Red Hat.

While generative AI holds immense potential, the costs associated with procuring, training, and fine-tuning large language models (LLMs) can be staggering, with some leading models costing nearly US$ 200 million to train. This doesn't include the expenses of aligning models with specific organisational data, which often requires specialised expertise.

Red Hat envisions that over the next decade, smaller, more efficient, and purpose-built AI models will become a significant part of the enterprise IT stack, alongside cloud-native applications. To make this future a reality, generative AI must become more accessible in terms of cost, contribution, and deployment across hybrid cloud environments. Red Hat believes that open source collaboration, which has historically solved complex software challenges, can similarly reduce the barriers to adopting generative AI.

RHEL AI is now generally available through the Red Hat Customer Portal for on-premises use or as a 'bring your own subscription' (BYOS) option on AWS and IBM Cloud. A BYOS offering for Azure and Google Cloud is planned for Q4 2024, and RHEL AI is also expected to be offered as a service on IBM Cloud later this year.

Hillery Hunter, CTO and general manager of innovation, IBM Infrastructure, said, "RHEL AI on IBM Cloud is bringing open source innovation to the forefront of gen AI adoption, allowing more organisations and individuals to access, scale and harness the power of AI. With RHEL AI bringing together the power of InstructLab and IBM's family of Granite models, we are creating gen AI models that will help clients drive real business impact across the enterprise."

Dell Technologies and Red Hat partner to bring RHEL AI to PowerEdge servers

Dell Technologies and Red Hat, Inc., have collaborated to bring Red Hat Enterprise Linux AI (RHEL AI) to Dell PowerEdge servers. RHEL AI, a foundation model platform built on an AI-optimised operating system, enables users to streamline the