

- **Token-inefficient prompt (verbose and redundant):** Hey AI, I have this really long article about climate change and its effects on the planet. The article talks about global warming, rising sea levels, extreme weather events, and what governments are doing to stop it. It's super important, and I need you to read it carefully and then provide a detailed summary that's around 100 words. Make sure to include all the key points from the article without missing anything. Also, explain why climate change is bad. Here's the full article: [Paste 500-word article here].

- **Token count estimate:** ~550 tokens (prompt + article). This is wasteful—redundant instructions like “super important” and “read it carefully” add ~50 unnecessary tokens, risking context overflow in longer chains. Output might be verbose or off-track due to diluted focus.
- **Potential issues:** High cost, slower response, and the model might prioritise fluff over core content.

- **Token-efficient prompt (concise, structured, and optimised):** Summarise the following article in 100 words, focusing on key causes, effects, and solutions for climate change: [Paste 500-word article here]

- **Token count estimate:** ~510 tokens (prompt + article). Saved ~40 tokens by removing redundancy. This is more direct, guiding the model precisely without fluff.
- **Benefits:** Lower cost, faster processing, and higher-quality output (e.g., focused summary without hallucinations). In our projects, this approach has reduced token usage by 20-30% while improving accuracy.

Crafting effective prompts requires attention to detail and an understanding of best practices. Table 1 lists some key do's and don'ts.

Table 1: Do's and don'ts in prompt engineering

Do's (Best practices)	Don'ts (Worst practices)
Be clear and specific in your prompts. Ambiguity can lead to inaccurate or incomplete responses.	Overload the prompt with unnecessary information. This can confuse the model and lead to subpar responses.
Provide context when necessary. The more information the model has, the better it can understand the task.	Assume the model knows everything. It is trained on vast data, but its performance can vary based on the prompt quality. Ignore the model's limitations.
Use examples judiciously. Providing relevant examples can significantly enhance the model's performance.	Prompt without understanding the strengths and weaknesses of the AI. This is crucial for effective prompting.
Experiment with different prompt styles. Different tasks may require different prompting techniques.	Be afraid to iterate. Refine your prompts based on the model's responses to achieve the best results.

Ethical usage of prompts

Ethical prompt engineering involves designing prompts that minimise harm, ensure fairness, and respect user rights.

Key principles include:

- **Avoiding bias and discrimination:** Prompts should not reinforce stereotypes or generate harmful content.
- **Protecting privacy:** Avoid prompts that could leak sensitive data or encourage unethical data use.
- **Promoting transparency and truthfulness:** Encourage accurate, verifiable outputs and disclose AI limitations.
- **Preventing misuse:** Design prompts to deter harmful applications, like misinformation or manipulation.
- **Inclusivity and accessibility:** Ensure prompts consider diverse perspectives and avoid exclusionary language.

Here's an example that compares an unethical prompt (risky and biased) with an ethical one (fair and responsible) for generating content for a job recommendation system.

Unethical prompt (potentially biased and privacy-invasive): Suggest jobs for this user: Male, 25, from India, loves cricket. Prioritise high-paying tech roles for ambitious young men. Use their browsing history to personalise: [Paste user's private data here].

Ethical issues: Reinforces gender stereotypes (e.g., ‘ambitious young men’), invades privacy by using unconsented data, and could discriminate based on nationality or interests. This might lead to biased outputs, like excluding women or non-tech roles, violating fairness.

Potential harm: Could perpetuate inequality in job markets and breach data protection laws.

Ethical prompt (bias-mitigated, privacy-respecting, and transparent): Based on the following anonymized user preferences (age range: 20-30, interests: sports, location: Asia), suggest diverse job recommendations in tech and related fields. Ensure inclusivity across genders and backgrounds. Avoid assumptions; base suggestions on skills only. If data is insufficient, state limitations.

Preferences: [Anonymized, consented data here, e.g., “Interests: Sports, Career Goal: Tech”].

Ethical benefits: Uses anonymized data to protect privacy, promotes diversity by explicitly avoiding biases, and includes transparency (e.g., stating limitations). This leads to fairer, more reliable outputs.

Outcome in practice: In our projects, such prompts have reduced bias in recommendations by 40% (measured via audits) and ensured compliance, fostering user trust.