Challenges and trade-offs in observability

Observability gives teams a clear window into complex systems. But that clarity comes with challenges. The first and most common is data overload. Every part of a system, such as servers, services and applications, can produce logs, metrics and traces. In a microservices setup with dozens or even hundreds of services, this can mean millions of data points every hour. Even with open source tools, storing and processing all this information can be expensive. For example, if a company records every API request and trace detail using Prometheus and Loki, the storage costs for just a few weeks of data can grow quickly.

Another challenge is tool fragmentation. Many teams use one tool for metrics, another for logs and a third for traces. One may be using Prometheus to track performance, Fluentd or Loki to handle log data, and Jaeger may follow request flows. When these tools do not integrate well, it becomes harder to follow a problem from start to finish. This slows down troubleshooting and can frustrate engineers. To address this, many teams are turning to OpenTelemetry, as it collects metrics, logs and traces in a single unified way.

Cost management is another trade-off. The free nature of open source tools does not translate to free operation at large scale. Servers, storage and network resources all add up. Teams manage costs by filtering out unneeded logs, sampling portions of data, and setting up storage duration restrictions.

Finally, there is a skills gap. Not every engineer is comfortable reading trace graphs, writing complex queries or building custom dashboards. The solution to this challenge demands training programmes, a simple dashboard design for all team members, and user-friendly observability tools.

When managed well, observability remains a powerful asset rather than an overwhelming burden.

The future: AIOps and automated remediation

Observability tools now serve beyond their traditional role of detecting problems. They predict system issues early and start to fix them before they cause operational disruptions thanks to AIOps technology (artificial intelligence for IT operations). Observability tools collect vast amounts of data, which AIOps processes through machine learning algorithms to identify potential problems.

AIOps systems detect patterns of memory usage reaching specific levels to anticipate service crashes before failures occur. The system detects problems ahead of time so it can both notify users and automatically restart services. Over time, it learns what ‘normal’ looks like for your system and quickly flags anything unusual, sometimes before users even notice.

Open source developers are currently exploring this area. The Keptn project develops automated remediation workflows for Kubernetes, which enables systems to initiate responses to failures without human involvement. Grafana plugins now include AI-powered anomaly detection capabilities that enhance both speed and accuracy of problem identification while reducing false alarm occurrences. The current developments point towards a future system that automates diagnosis and real-time responses.

However, technology is only half the story. True observability depends just as much on team culture as it does on tools. The system functions optimally when developers work alongside testers and operations staff to share responsibilities in maintaining system health. Writing better logs, asking smarter questions, and learning from past incidents should be part of everyday work.

Observability functions as the DevOps nervous system because it continuously monitors system activities to detect threats and trigger immediate efficient responses. As open source tools and AIOps systems become more sophisticated, teams that integrate observability principles throughout their development process will be the most successful. END

By: Dr Chinmoy Kumar

The author is an associate professor in business analytics at SCMS - PG, Dayananda Sagar University, Bengaluru. His research and teaching interests include analytics, data science, emerging technologies, business models and e-governance for sustainable development. He is also the founder of ‘The Learning Network’, a pro bono educational venture based in Ranchi, Jharkhand, India.